

Scikit-learn and Tour of Classifiers


Support Vector Machines



Support Vector Machines


- Widely-used, powerful classifiers
- Can be considered an extension of the Perceptron
- Support vector classification (SCV) objective: **maximize margin**
- Margin: the distance between the separating hyperplane (decision boundary) and the training samples that are closest to this hyperplane
- The training samples closest to the hyperplane are called **support vectors**
- Kernel-SVM for nonlinear decision boundaries, resulting predictor is sparse

Support Vector Machines



Which hyperplane?

Margin (gap between decision boundary and hyperplanes)



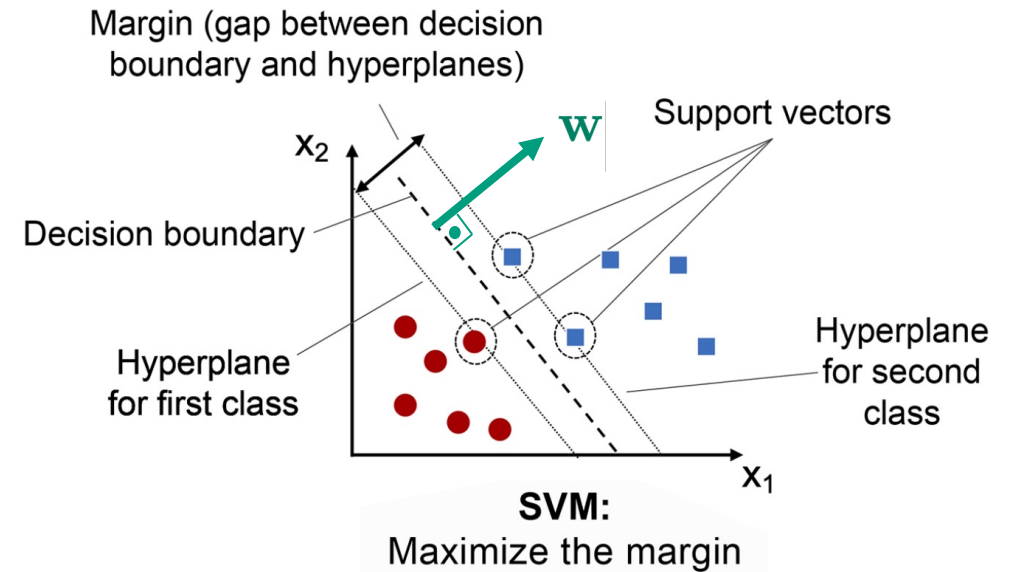
SVM:
Maximize the margin

Maximum margin classifier

- Decision boundaries with **large margins** tend to have a **lower generalization error**
- Decision boundaries with **small margins** are more **prone to overfitting**
- **Linear decision boundary:**

$$\mathbf{x}^{(i)} \boldsymbol{\theta} = b + \mathbf{x}^{(i)} \mathbf{w} = 0$$

normal vector
on hyperplane



$$\mathbf{x}^{(i)} = \begin{bmatrix} 1 & x_1^{(i)} & x_2^{(i)} & \dots & x_m^{(i)} \end{bmatrix} \quad \boldsymbol{\theta} = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_m \end{bmatrix} := \begin{bmatrix} b \\ \mathbf{w} \end{bmatrix} \quad \mathbf{w} = \begin{bmatrix} \theta_1 \\ \theta_2 \\ \vdots \\ \theta_m \end{bmatrix}, \quad b = \theta_0$$

- **Distance** between positive and negative hyperplanes: $2 / \|\mathbf{w}\|$
- **Maximize distance** while **classifying samples correctly**

Maximum margin (linearly separable)

Margin (gap between decision boundary and hyperplanes)


- Derivation of maximum margin formula not syllabus; see e.g. <https://www.cs.cornell.edu/courses/cs4780/2018fa/lectures/lecturenote09.html>
- Finding the **maximum margin** hyperplanes can be cast into the **quadratic optimization problem**:

$$\min_{\mathbf{w}, b} \mathbf{w}^T \mathbf{w}$$

$$\text{such that for all } i \quad y^{(i)} (\mathbf{x}^{(i)} \mathbf{w} + b) \geq 1$$

- For **support vectors** (data on the hyperplanes), it holds

$$y^{(i)} (\mathbf{x}^{(i)} \mathbf{w} + b) = 1$$





$$\text{Classes: } y^{(i)} \in \{-1, 1\}$$

Recall that: $\mathbf{w}^T \mathbf{w} = \|\mathbf{w}\|_2^2 = \sum_{i=1}^m w_i^2$

Maximum margin – soft margin

- Often, data is not linearly separable. Idea: soft margin by introducing “**slack variables**” $\xi^{(i)}$



Maximum margin – soft margin

- Often, data is not linearly separable. Idea: soft margin by introducing “**slack variables**” $\xi^{(i)}$
- Modified optimization problem:

$$\min_{\mathbf{w}, b} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^n \xi^{(i)} \quad \text{such that for all } i \quad y^{(i)} (\mathbf{x}^{(i)} \mathbf{w} + b) \geq 1 - \xi^{(i)} \quad \text{and} \quad \xi^{(i)} \geq 0$$

- The slack variables allow for some samples to be within the margin (or even on the wrong side = misclassified)
- For high **C** we penalize sample in or on the wrong side of the margin strongly
- The optimization problem with slack variables can be written in terms of the loss function:



$$L(\boldsymbol{\theta}) = L(\mathbf{w}, b) = \underbrace{\mathbf{w}^T \mathbf{w}}_{\ell_2\text{-regul.}} + C \sum_{i=1}^n \underbrace{\max \left\{ 1 - y^{(i)} (\mathbf{x}^{(i)} \mathbf{w} + b), 0 \right\}}_{\text{hinge-loss}}$$

Recall that:

$$\mathbf{w}^T \mathbf{w} = \|\mathbf{w}\|_2^2 = \sum_{i=1}^m w_i^2$$


Maximum margin – soft margin

- **Effect of C:** decreasing C increases the bias (underfitting) and lowers the variance (overfitting) of the model



SVM code example


03_svm_linear_blobs.ipynb




SVM code example

03_svm_linear_iris.ipynb


SVC with linear kernel, $C=0.01$

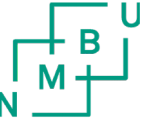


SVC with linear kernel, $C=1.0$



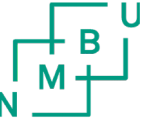
SVC with linear kernel, $C=1000.0$





Logistic regression versus Linear SVM

- LR and linear SVM often yield very similar results
- **Logistic regression** tries to **maximize** the **conditional likelihoods** for all of the training data, which makes it **more prone** to **outliers** than **SVMs**
- **SVMs** mostly care about the samples that are **closest** to the **decision boundary** (**support vectors**)
- **Logistic regression** has the advantage that it is a **simpler model** and can be implemented more easily



Linear model (linear decision boundaries)

- **Linear model classifiers** may **appear** to be very **restrictive** in **low-dimensional** spaces (**very few** features in X) because of their **straight line or plane** boundaries
- For **high dimensional** data (**many** features in X) linear model classifiers may act as a **guard against** overfitting



Scikit-learn implementation

- `Perceptron`, `LogisticRegression` and `LinearSVC` classes in scikit-learn
make use of the **LIBLINEAR** library (highly optimized C/C++ library developed at the National Taiwan University)
- `SVC` and `SVM` classes in scikit-learn
makes use of **LIBSVM** library (an equivalent C/C++ library specialized for SVMs developed at the National Taiwan University)
- LIBLINEAR and LIBSVM are **faster** than **native** Python implementations



Scikit-learn implementation


- `Perceptron`, `LogisticRegression` and `LinearSVC` classes in scikit-learn
make use of the **LIBLINEAR** library (highly optimized C/C++ library developed at the National Taiwan University)
- `SVC` and `SVM` classes in scikit-learn
makes use of **LIBSVM** library (an equivalent C/C++ library specialized for SVMs developed at the National Taiwan University)
- LIBLINEAR and LIBSVM are **faster** than **native** Python implementations



Scikit-learn alternative implementation

- When datasets are **too large** to fit into computer memory
 - Alternative implementations using the `SGDClassifier` class
 - `SGDClassifier` class supports online learning via the `partial_fit` method
 - Concept behind the `SGDClassifier` class is **the stochastic gradient descend algorithm**
- **Initialize** the stochastic gradient descent version of the **perceptron (ppn)**, **logistic regression (lr)**, and a **support vector machine (svm)** with default parameters as follows

```
from sklearn.linear_model import SGDClassifier
ppn = SGDClassifier(loss='perceptron')
lr = SGDClassifier(loss='log')
svm = SGDClassifier(loss='hinge')
```




Nonlinear decision boundaries

Kernel Support Vector Machines

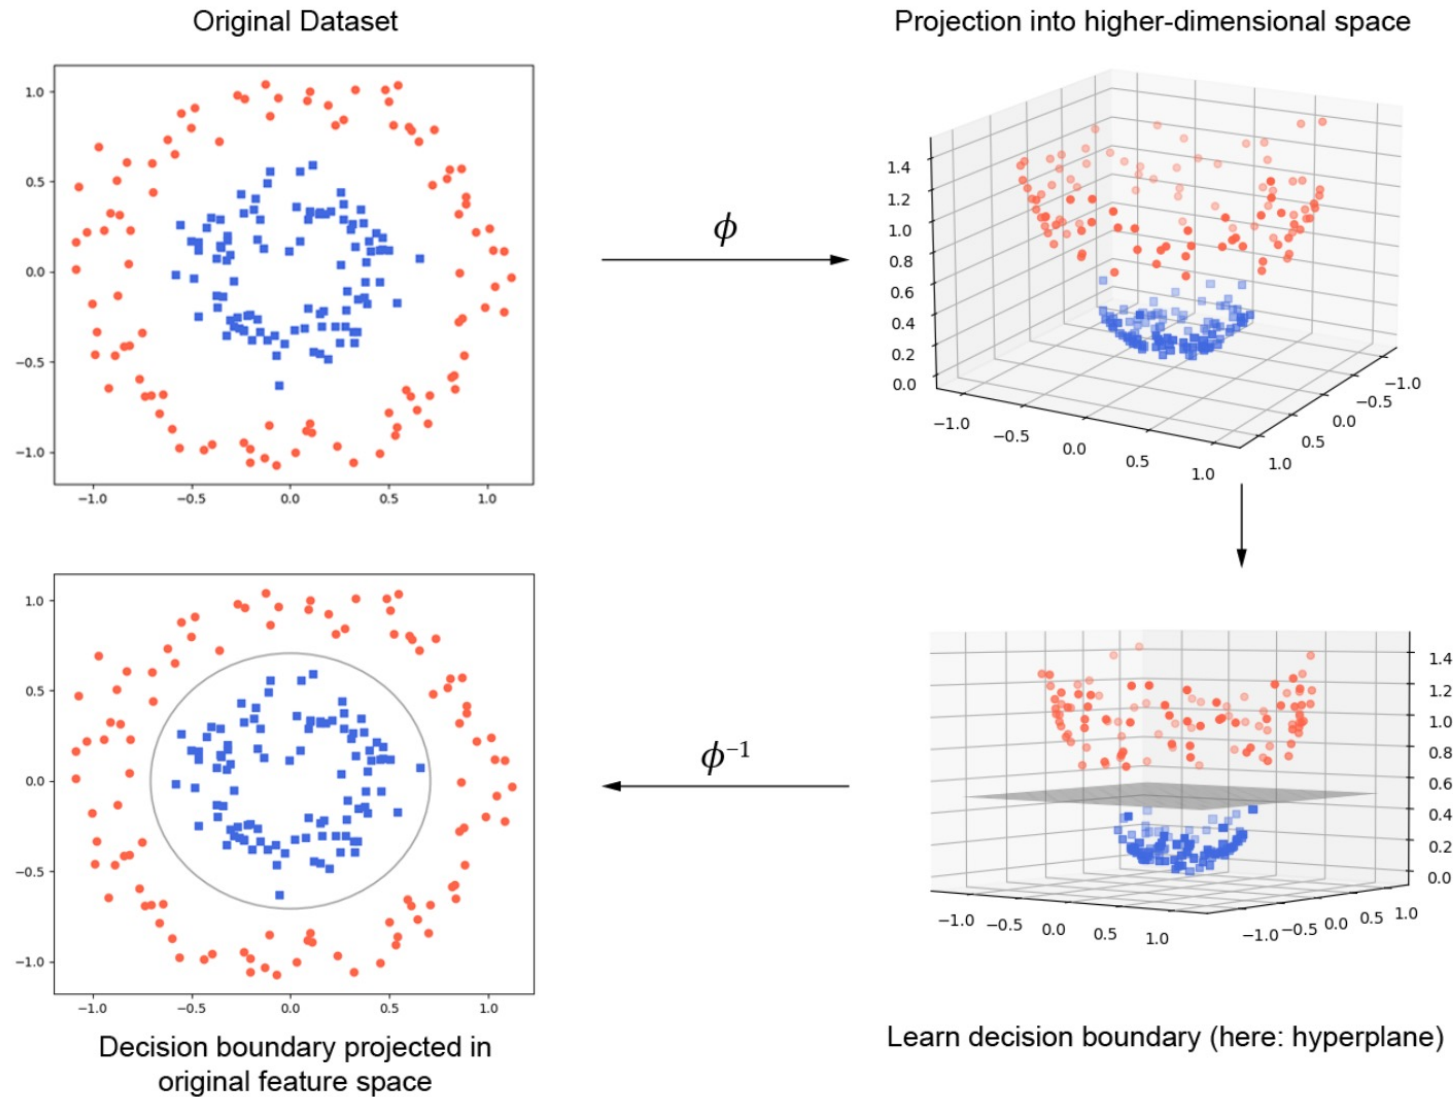
Nonlinear decision boundaries – Motivation

Data not linearly separable!



“XOR” dataset

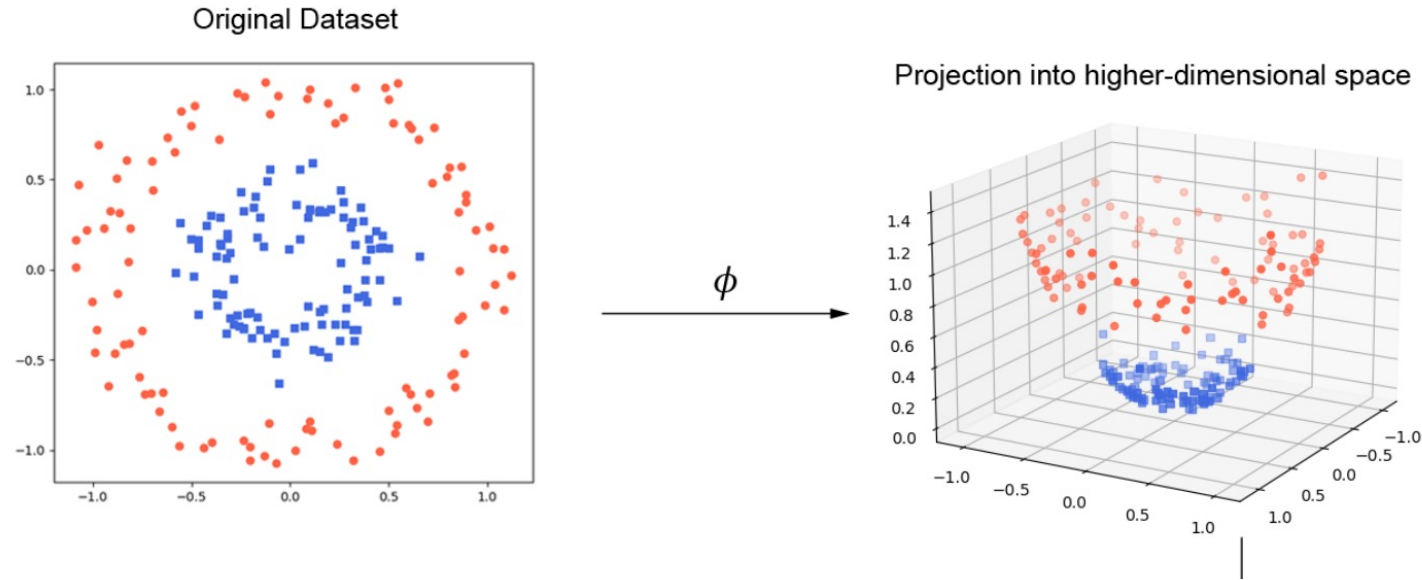
Idea behind kernel methods



ϕ feature mapping

Data may be linearly separable in a higher dimensional feature space

Idea behind kernel methods



ϕ feature mapping, here:

$$\phi(x_1, x_2) = (z_1, z_2, z_3) = (x_1, x_2, x_1^2 + x_2^2)$$

Kernel trick

sample: $\mathbf{x}^{(i)} = \begin{bmatrix} 1 & x_1^{(i)} & x_2^{(i)} & \dots & x_m^{(i)} \end{bmatrix}$



- Many algorithms can be formulated in terms of **inner products** of the **samples**

$$\mathbf{x}^{(i)} \mathbf{x}^{(j)T} = \sum_{k=0}^m x_k^{(i)} x_k^{(j)}$$

- After **feature mapping**, we have an **inner product** in a **high-dimensional space**

$$\phi(\mathbf{x}^{(i)}) \phi(\mathbf{x}^{(j)})^T$$

- Often, there is an **easier way** to compute the resulting number in the original space

$$\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \phi(\mathbf{x}^{(i)}) \phi(\mathbf{x}^{(j)})^T$$

Kernel: $\kappa(\mathbf{x}, \mathbf{x}') : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$

without ever computing $\phi(\mathbf{x}^{(i)})$

- $\phi(\mathbf{x}^{(i)})$ is never needed. There is even kernels for which no explicit finite dimensional feature mapping exists

Kernel trick

- **Example: A quadratic polynomial kernel, $m=3$ (number of features)**

$$\begin{aligned}
 \kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) &= \left(\mathbf{x}^{(i)T} \mathbf{x}^{(j)} \right)^2 \\
 &= \left(\sum_{k=1}^m x_k^{(i)} x_k^{(j)} \right) \left(\sum_{k=1}^m x_k^{(i)} x_k^{(j)} \right) \\
 &= \sum_{k=1}^m \sum_{l=1}^m x_k^{(i)} x_l^{(i)} x_k^{(j)} x_l^{(j)} \\
 &= \sum_{k=1}^m \sum_{l=1}^m (x_k^{(i)} x_l^{(i)}) (x_k^{(j)} x_l^{(j)})
 \end{aligned}$$

Computation in original
space via kernel: $O(m)$

Therefore, corresponding feature mapping is:

$$\phi(\mathbf{x}^{(i)}) = \left[x_1^{(i)} x_1^{(i)}, x_1^{(i)} x_2^{(i)}, x_1^{(i)} x_3^{(i)}, x_2^{(i)} x_1^{(i)}, x_2^{(i)} x_2^{(i)}, x_2^{(i)} x_3^{(i)}, x_3^{(i)} x_1^{(i)}, x_3^{(i)} x_2^{(i)}, x_3^{(i)} x_3^{(i)} \right]$$

$$\phi(\mathbf{x}^{(i)}) \phi(\mathbf{x}^{(j)})^T = \sum_{p=1}^3 \phi(\mathbf{x}^{(i)})_p \phi(\mathbf{x}^{(j)})_p = \sum_{k=1}^{m=3} \sum_{l=1}^{m=3} (x_k^{(i)} x_l^{(i)}) (x_k^{(j)} x_l^{(j)})$$

Computation in feature
space: $O(m^2)$



Kernel trick

- If the data only occurs in form of **inner products** in an ML algorithm, we can replace these by evaluations of a **kernel function** to compute inner products in a high-dimensional feature space
- Evaluation with data in the original space is cheaper
- Evaluation in the feature space is expensive
- We might not even know what the feature mapping looks like!

Kernel functions

- Kernel can be seen to measure **similarity** between samples
- “**Valid**” kernel (**Mercer** kernel): if kernel matrix is **positive semi-definite**

$$K_{ij} = \kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \begin{bmatrix} \kappa(\mathbf{x}^{(1)}, \mathbf{x}^{(1)}) & \cdots & \kappa(\mathbf{x}^{(1)}, \mathbf{x}^{(n)}) \\ \vdots & \ddots & \vdots \\ \kappa(\mathbf{x}^{(n)}, \mathbf{x}^{(1)}) & \cdots & \kappa(\mathbf{x}^{(n)}, \mathbf{x}^{(n)}) \end{bmatrix} \quad \kappa(\mathbf{x}, \mathbf{x}') : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$$

a matrix \mathbf{A} is positive semi-definite if $\mathbf{v}^T \mathbf{A} \mathbf{v} \geq 0$ for all non-zero \mathbf{v}

- For valid kernels, it is guaranteed that a corresponding feature mapping exists (although, it might be infinite dimensional)

Typical kernel functions

- Linear kernel

$$\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \phi(\mathbf{x}^{(i)})\phi(\mathbf{x}^{(j)})^T = \mathbf{x}^{(i)}\mathbf{x}^{(j)T}$$

- Polynomial kernel (tuning parameter p)

$$\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = (\mathbf{x}^{(i)}\mathbf{x}^{(j)T} + 1)^p, \quad \text{for polynomial degree } p$$

- Isotropic Gaussian kernel / Radial basis function (RBF) kernel (tuning parameter γ)

$$\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \exp\left(-\frac{\|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|_2^2}{2\sigma^2}\right) := \exp\left(-\gamma\|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|_2^2\right)$$



More kernel functions (not syllabus)


- There is also kernel functions for strings, text, etc.
- Kernelized algorithm can then operator directly on that data without explicit transformation
- Choice of **kernel function** should be **tailored to the data**



Kernel SVM

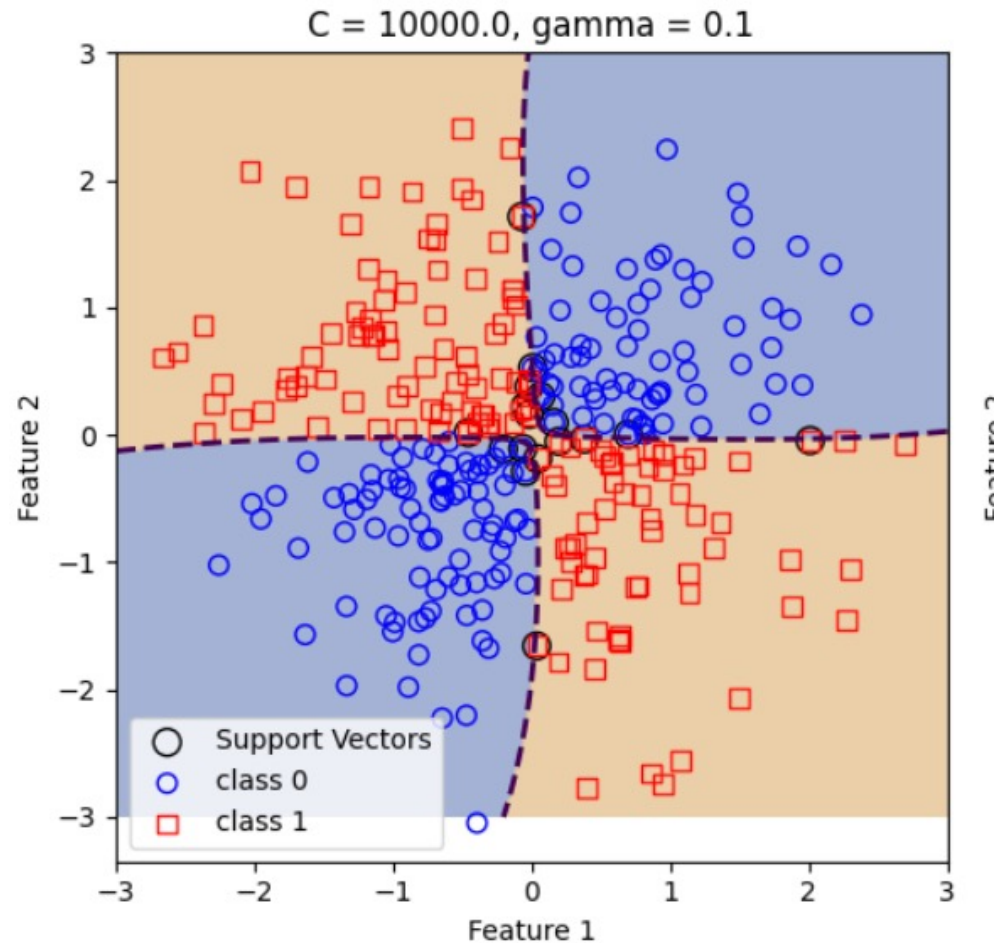
- In prediction and optimization algorithm (we didn't look at it in detail) data only occurs as **inner products** → **replace inner products** by **kernel function** calls
- In scikit-learn the type of kernel is a parameter of the SVC (Support Vector Classifier)
- The kernels have hyperparameters (e.g. `degree` for polynomial kernel)

Kernel SVM Code example



“XOR” dataset

Kernel SVM Code example (RBF kernel)



Chosen kernel function:

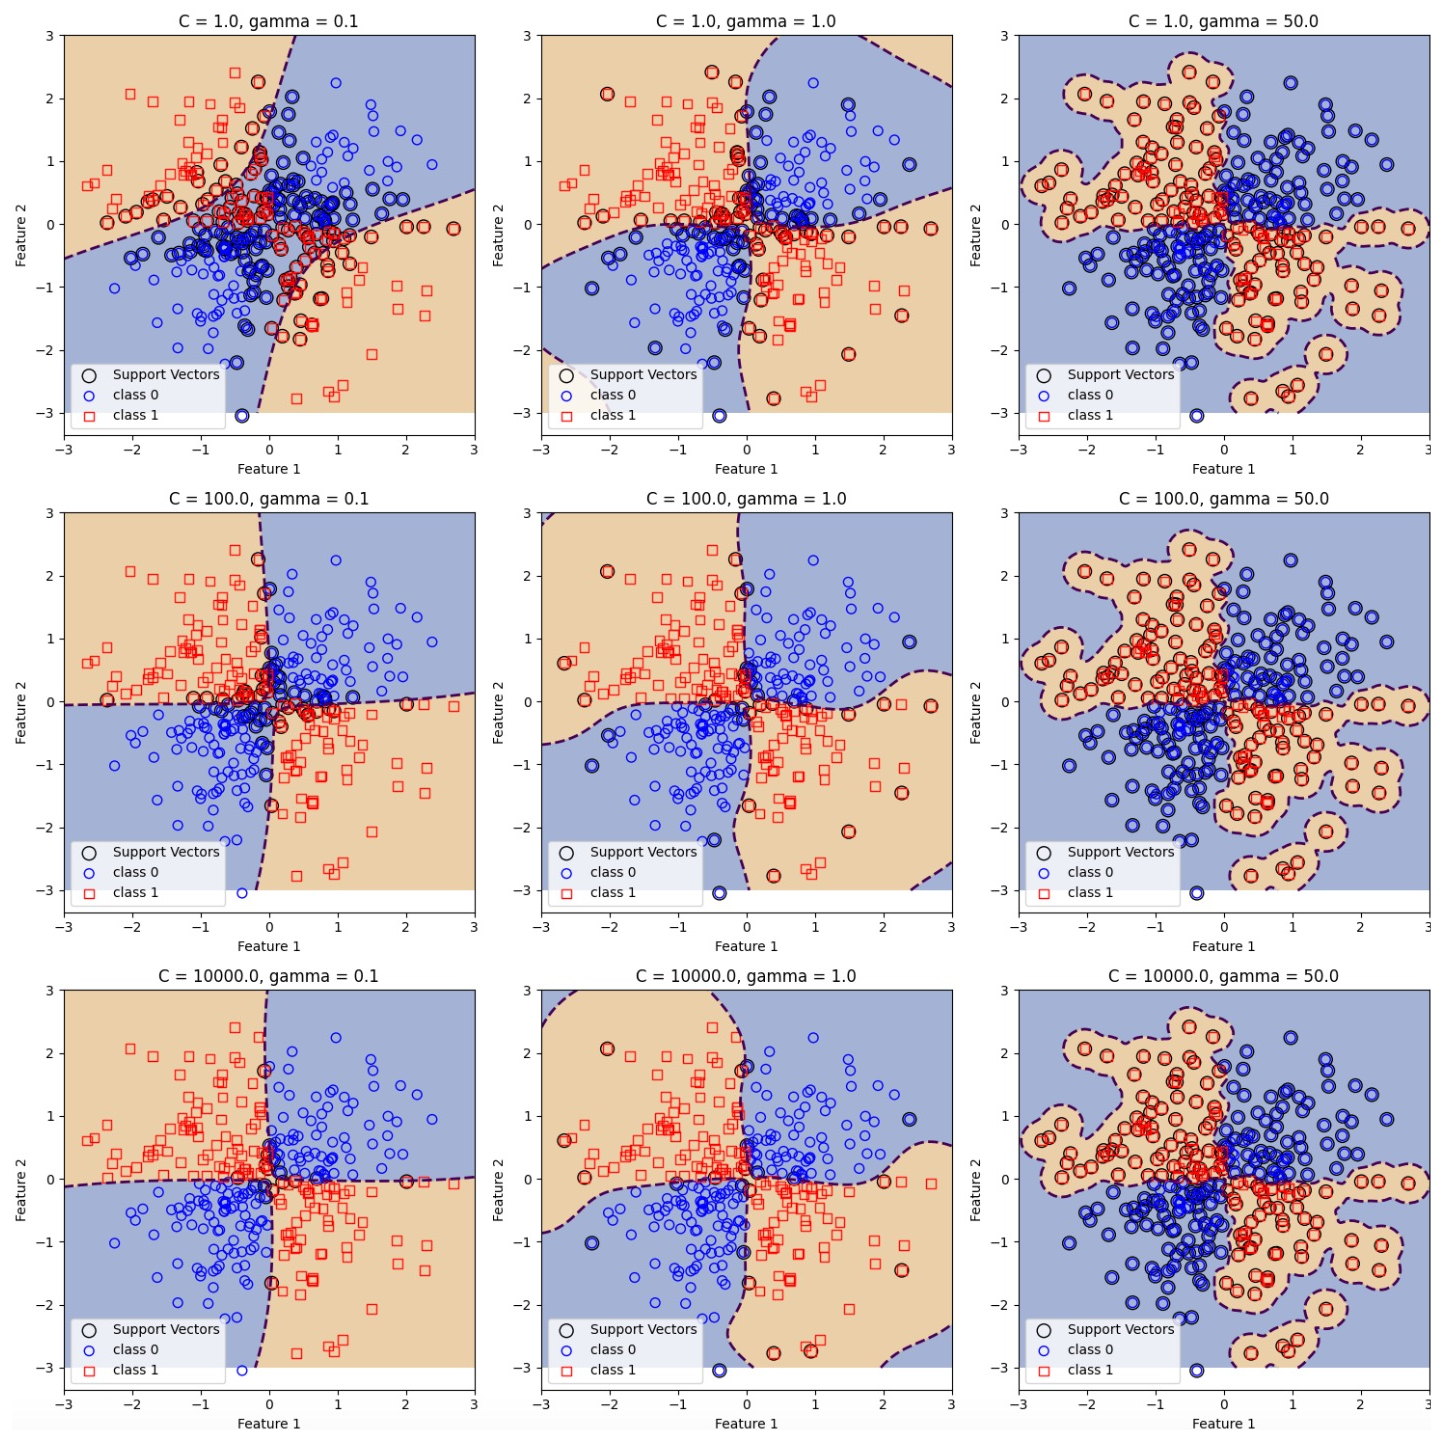
$$\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) := \exp \left(-\gamma \|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|_2^2 \right)$$

Kernel SVM

- We need to tune both C and the kernel parameter(s)
- For example: grid search
 - C in $[1.0, 100.0, 10000.0]$
 - γ in $[0.1, 1.0, 50.0]$

`03_svm_kernel_xor.ipynb`

`03_svm_gridsearch_iris.ipynb`



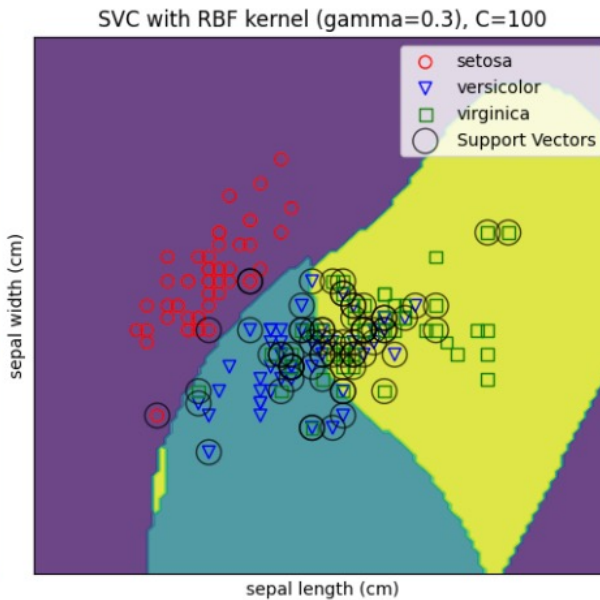
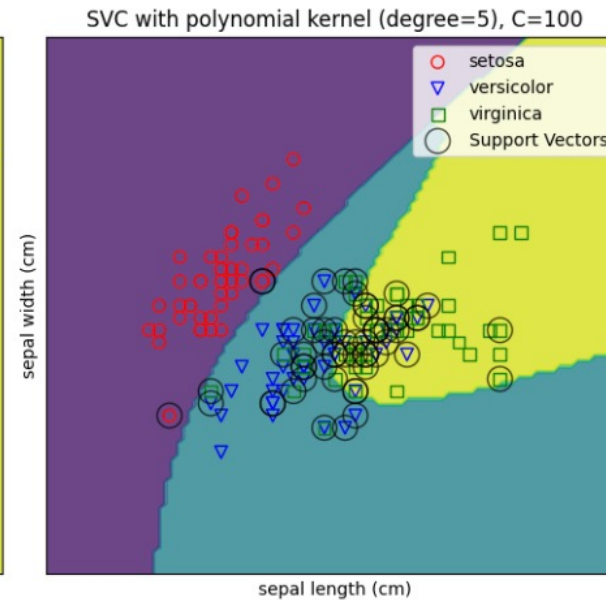
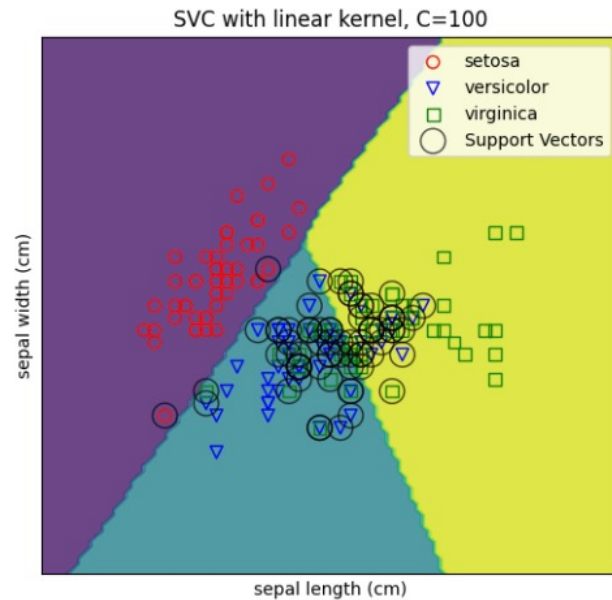
Kernel SVM

- We can also try different kernels
- For example:
 - Linear
 - Polynomial
 - RBF

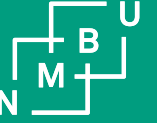
$$\mathbf{x}^{(i)} \mathbf{x}^{(j)T}$$

$$(\mathbf{x}^{(i)} \mathbf{x}^{(j)T} + 1)^p$$

$$\exp(-\gamma \|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|_2^2)$$



03_svm_different_kernels_iris.ipynb




Other kernelized ML algorithms

Nonlinear decision boundaries



Kernelized ML algorithms

- Many of the presented algorithms can be **kernelized** by replacing inner products with kernel evaluations (sometimes after reformulation of the original algorithm)
- Examples:
 - **Kernel Perceptron**
 - **Kernel Logistic Regression**
 - **Kernel Ridge Regression** (linear regression + L2 regularizer)
 - **Kernel SVM**
 - **Kernel k-nearest neighbors** (next lecture)



Nonlinear decision boundaries

Kernel perceptron

Kernel perceptron

- To gain intuition of how to kernelize linear models
- Recall that for the perceptron $y^{(i)} \in \{1, -1\}$


– we **predict** as

$$\hat{y}^{(i)} = \text{sign}\{\mathbf{x}^{(i)}\boldsymbol{\theta}\}$$

– we **train** as

```

 $\boldsymbol{\theta} \leftarrow \mathbf{0}$ 
for  $i = 1, \dots, n$  do
     $\hat{y}^{(i)} \leftarrow \text{sign}\{\mathbf{x}^{(i)}\boldsymbol{\theta}\}$ 
    if  $\hat{y} \neq y^{(i)}$  then
         $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \eta y^{(i)} \mathbf{x}^{(i)T}$ 
    end if
end for
    
```



this actually means: $\boldsymbol{\theta} = \sum_{j=1}^n \alpha^{(j)} \eta y^{(j)} \mathbf{x}^{(j)T}$

$\alpha^{(j)}$: number of times sample j is misclassified

Kernel perceptron

- Parameters are linear combination of the samples: $\theta = \sum_{j=1}^n \alpha^{(j)} \eta y^{(j)} \mathbf{x}^{(j)T}$
- Recall that for the perceptron $y^{(i)} \in \{1, -1\}$

– we **predict** as

$$\hat{y}^{(i)} = \text{sign}\{\mathbf{x}^{(i)} \theta\}$$

– we **train** as

```

 $\theta \leftarrow 0$ 
for  $i = 1, \dots, n$  do
   $\hat{y}^{(i)} \leftarrow \text{sign}\{\mathbf{x}^{(i)} \theta\}$ 
  if  $\hat{y} \neq y^{(i)}$  then
     $\theta \leftarrow \theta + \eta y^{(i)} \mathbf{x}^{(i)T}$ 
  end if
end for

```

“dual” formulation

$$\hat{y}^{(i)} = \text{sign} \left\{ \sum_{j=1}^n \alpha^{(j)} \eta y^{(j)} \mathbf{x}^{(i)} \mathbf{x}^{(j)T} \right\}$$

$\alpha \leftarrow 0$

for $i = 1, \dots, n$ **do**

$$\hat{y}^{(i)} \leftarrow \text{sign} \left\{ \sum_{j=1}^n \alpha^{(j)} \eta y^{(j)} \mathbf{x}^{(i)} \mathbf{x}^{(j)T} \right\}$$


if $\hat{y} \neq y^{(i)}$ **then**

$$\alpha^{(i)} \leftarrow \alpha^{(i)} + 1$$

end if

end for

inner product



Kernel perceptron

- Parameters are linear combination of the samples: $\theta = \sum_{j=1}^n \alpha^{(j)} \eta y^{(j)} \mathbf{x}^{(j)T}$
- Dual formulation $y^{(i)} \in \{1, -1\}$

– we **predict** as

$$\hat{y}^{(i)} = \text{sign} \left\{ \sum_{j=1}^n \alpha^{(j)} \eta y^{(j)} \mathbf{x}^{(i)} \mathbf{x}^{(j)T} \right\}$$

– we **train** as

```

 $\alpha \leftarrow 0$ 
for  $i = 1, \dots, n$  do
     $\hat{y}^{(i)} \leftarrow \text{sign} \left\{ \sum_{j=1}^n \alpha^{(j)} \eta y^{(j)} \mathbf{x}^{(i)} \mathbf{x}^{(j)T} \right\}$ 
    if  $\hat{y} \neq y^{(i)}$  then
         $\alpha^{(i)} \leftarrow \alpha^{(i)} + 1$ 
    end if
end for
    
```



Kernel perceptron

- Kernel perceptron

- we **predict** as

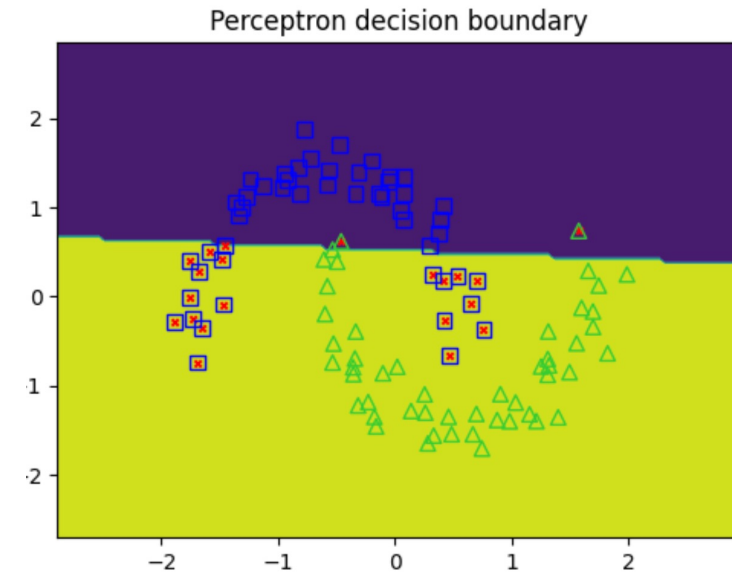
$$\hat{y}^{(i)} = \text{sign} \left\{ \sum_{j=1}^n \alpha^{(j)} y^{(j)} \kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) \right\}$$

- we **train** as

```

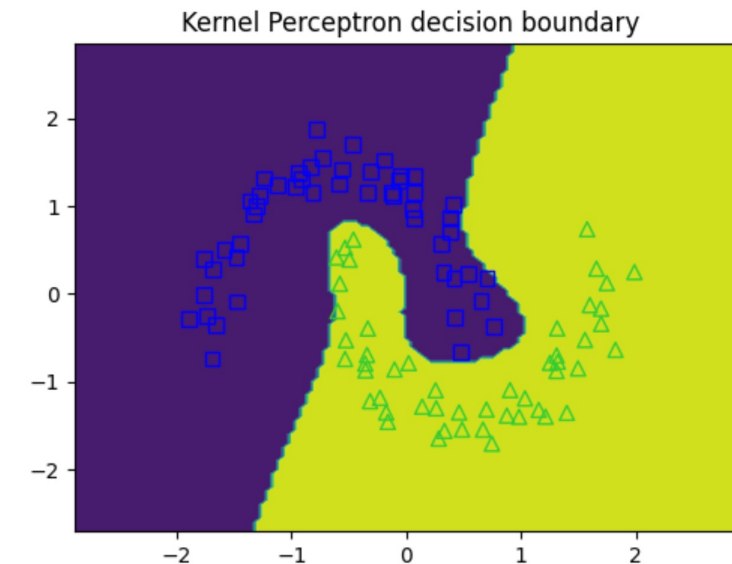
 $\alpha \leftarrow 0$ 
for  $i = 1, \dots, n$  do
     $\hat{y}^{(i)} \leftarrow \text{sign} \left\{ \sum_{j=1}^n \alpha^{(j)} y^{(j)} \kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) \right\}$ 
    if  $\hat{y} \neq y^{(i)}$  then
         $\alpha^{(i)} \leftarrow \alpha^{(i)} + 1$ 
    end if
end for

```

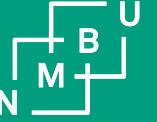


$(p=1)$

$$(\mathbf{x}^{(i)} \mathbf{x}^{(j)T} + 1)^p$$



$p=5$



Nonlinear decision boundaries

Kernel logistic regression (extra, not syllabus)

Kernel logistic regression

- Logistic regression can be **kernelized** analogously to the perceptron by switching to a dual formulation
- Generalized linear models (like logistic regression) can be turned into **kernel machines** by considering a feature mapping of the form:

$$\phi(\mathbf{x}^{(i)}) = [\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(i)}), \dots, \kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(n)})]$$

Number of features is number of samples (or somewhat less for sparse kernel machines)

- This corresponds to the realization that the weights can be written as a linear combination of the samples
- When only a subset of the samples is used (**support vectors**), we call the resulting model a **sparse kernel machine**



Kernel logistic regression versus Kernel SVM

- **Kernel SVMs** yields **sparse** kernel machines which are faster in prediction and have lower memory consumption
- **Kernelized Logistic regression** uses all data points in kernel matrix, resulting in a dense kernel machine
- **Kernelized Logistic regression** with sparsity-promoting regularization like **L1-regularization** yields a **sparse** predictor (samples with zero weights can be dropped)

