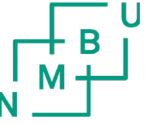


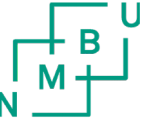
# Scikit-learn and Tour of Classifiers

## K-Nearest Neighbors



# K-Nearest Neighbors

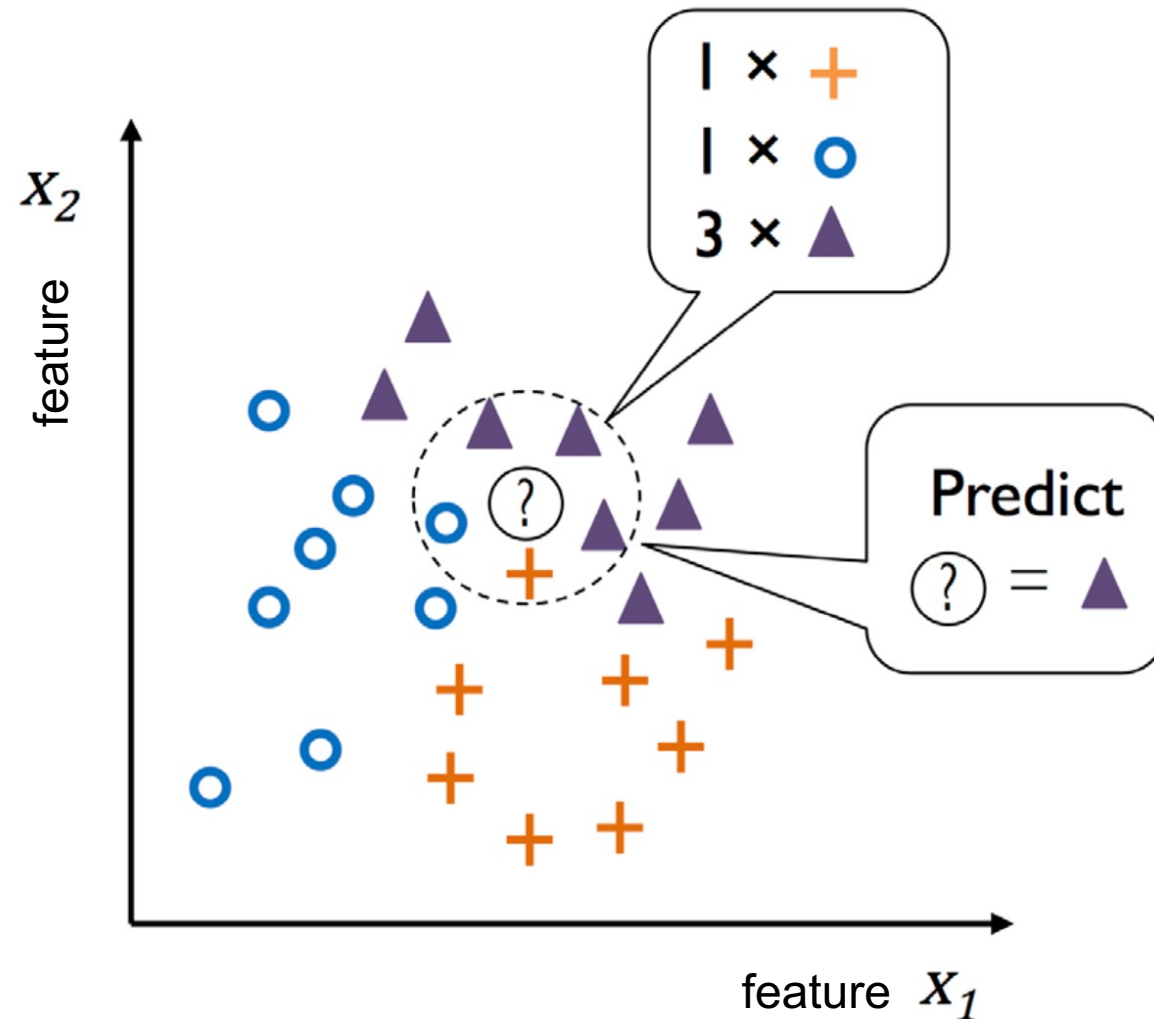
- Fundamentally different algorithm from what we learned so far
- Does not learn a discriminative function
- Memorizes the training data set and makes prediction based on that



# K-Nearest Neighbors – Algorithm

1. Choose parameter "**k**" and a **distance metric**
2. Find the k-nearest neighbors of a data record that we want to classify
3. Assign the class label by majority vote
  - *Tie break*: take the class of the closest neighbor and if still tied use lowest label

# K-Nearest Neighbors – Algorithm



**Example:** 5-nearest neighbors of a new data record (?) that we want to classify

# K-Nearest Neighbors – Distance metric

- A **distance metric** measures **distance** between two samples
- **Should fit the type of data**
- Common for real numbers: Euclidean distance

$$d(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \sqrt[p]{\sum_k |x_k^{(i)} - x_k^{(j)}|^p}$$

Minkowski distance (used in scikit-learn K-NN)

$p = 1$

$p = 2$


$$\sum_k |x_k^{(i)} - x_k^{(j)}|$$


Manhattan distance

$$d(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \sqrt{\sum_k |x_k^{(i)} - x_k^{(j)}|^2}$$

Euclidean distance

- Euclidean distance
- Manhattan distance





# K-Nearest Neighbors

- Tuning **k** (and distance metric) is crucial to find a good balance between overfitting and underfitting
  - Lower **k** means that the model is more complex and has higher variance
  - Higher **k** means less variance (consulting more neighbors) but higher bias (far away neighbors might be consulted)
  - Higher **k** means that prediction becomes more costly




# K-Nearest Neighbors


03\_knn.ipynb

- Code example, testing accuracy on test/train set for different  $k$


1-Nearest-Neighbor decision boundary




3-Nearest-Neighbor decision boundary



9-Nearest-Neighbor decision boundary







# Parametric versus non-parametric models



# Parametric versus non-parametric models

- **Parametric models**
  - Learn **parametrized** discriminative functional (fixed set of parameters)
  - Predict/classify new points **without need for** the training data
  - Examples: Perceptron, Adaline, Linear Regression, Logistic Regression, (*linear* SVM)
- **Non-parametric models**
  - No fixed set of params, no. of parameter **changes with amount of training data**
  - Training data (at least a subset) is needed for prediction
  - Examples: Decision trees, most kernel machines (Kernel-SVM, Kernel-Perceptron, ...)



# Some advantages and disadvantages of a memory-based non-parametric approach like K-NN

- Advantages:
  - Low (or zero) cost adaption to new training samples
  - Good predictive vs computational performance for small to medium-sized data sets
- Disadvantages:
  - Storage and prediction cost grows with the number of samples (prediction cost can be lowered by using efficient data structures, e.g. k-d trees for K-NN)
  - Curse of dimensionality → for a high number of features, prone to overfitting, no other training sample may be informative due to increasingly sparse feature space population with increasing dimension (number of features)

