

# Scikit-learn and Tour of Classifiers

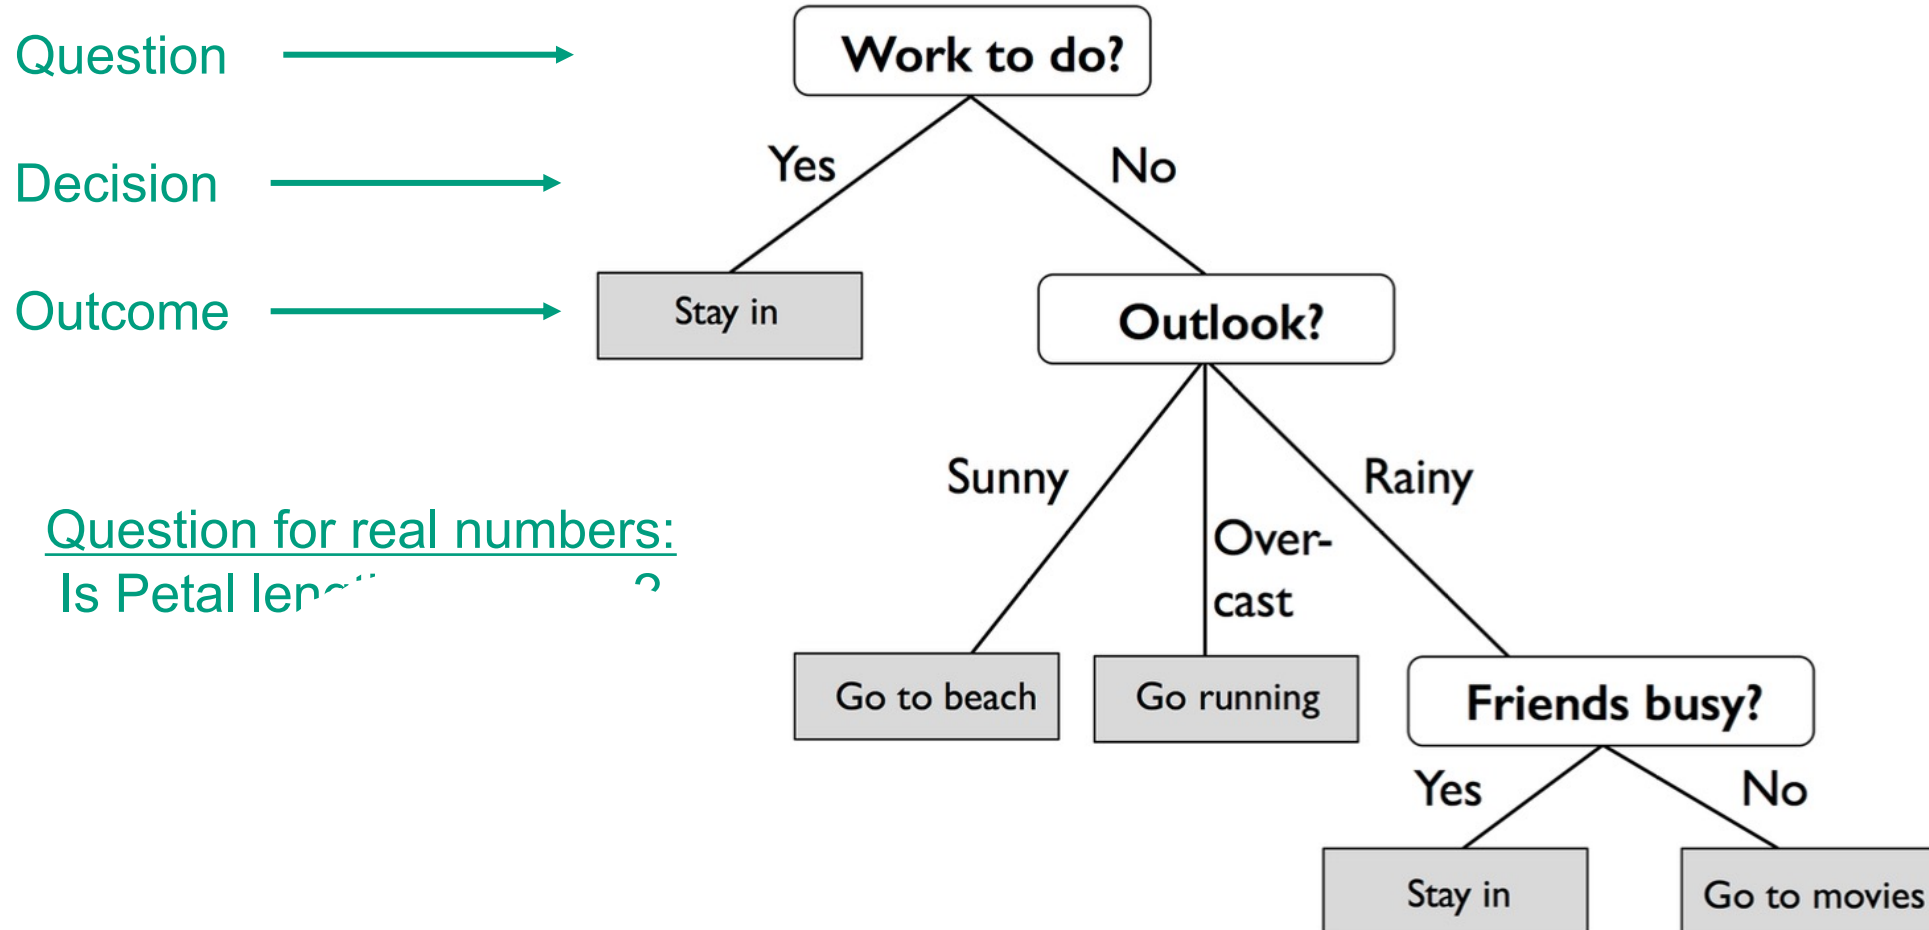
## Decision Tree learning



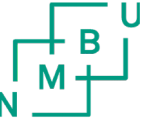
# Decision Tree learning

- Based on recursive binary **partition** of the **feature space**
- **Decide** class based on a series of “**questions**”
- Partition based on **axis-oriented** hyperplanes (“zero weights only bias”)
- The **number of partitions** is called **tree depth** and determines model complexity
- Very popular due to easy **interpretability** of the **resulting decision function**

# Decision trees – Illustrative example




Tree depth = 3



# Decision trees – Learning goal


- **Learn a series of questions** that will lead to a decision
- Which question? That is, which split do we choose?
  - Choose the split with the largest **information gain (IG)**
- We can repeat adding splits until **leaf nodes** are **pure** (all samples associated with the leaf node belong to the same class)
- Trees with only pure leaf nodes are usually too deep (overfitting!)
- **Pruning** a tree means removing branches or setting a limit for the maximum depth

# Decision trees – Example showing feature space





# Decision trees – Example showing feature space

**Pure leaf node**  
(here: “region”)  
Only contains  
“setosa” samples



Tree depth = 4





# Decision trees – Maximize information gain (IG)

- **Information gain (IG)** of splitting along feature “f” with threshold “t”, can be measured by

$$IG(D_p, (f, t)) = I(D_p) - \sum_{j=1}^{C=2} \frac{N_j(t)}{N_p} I(D_j(t))$$

- $D_n \subseteq D = \{X, y\}$  is the **subset** of the **data** at **node**  $n$  (“ $p$ ”: parent, “ $j$ ”:  $j$ -th child of parent  $p$ )
- $N_n$  is the number of samples in  $D_n$
- Typically,  $C=2$  (**number of child nodes**); yields efficient data structures (**binary trees**)
- Finally,  $I(D)$  is the **impurity** of the data
- Choose split along the **feature with the maximum information gain**  $\operatorname{argmax}_{\{f, t\}} IG(D_p, f, t)$



## Decision trees – Impurity measures (Classification error)

- **Impurity measure (I)** of a dataset associated with a node: a **measure** of how far away from a **pure** node (all samples one class) we are
- Let's give each node a unique index  $n$ ; and  $K$  is the number of classes
- $D_n \subseteq D = \{X, y\}$  is the subset of the data at node  $n$
- “Minimize probability of misclassification”

Classification error: 
$$I_E(D_n) = 1 - \max_{c \in K} \{ p(c | D_n) \}$$

$p(c | D_n)$  : Sample probability of a sample having class  $c$  in data subset  $D_n$

$$\rightarrow p(c | D_n) = \frac{N_c}{\sum_{c \in K} N_c} \qquad \sum_{c \in K} p(c | D_n) = 1$$




# Decision trees – Sample probability at node

$p(c \mid D_n)$  : Sample probability of a sample having class  $c$  in data subset  $D_n$

$$\rightarrow p(c \mid D_n) = \frac{N_c}{\sum_{c \in K} N_c}$$

$$\sum_{c \in K} p(c \mid D_n) = 1$$



$$p(\bullet \mid D_p) = 5/11 = 0.45$$

$$p(\blacksquare \mid D_p) = 6/11 = 0.55 = 1 - p(\bullet \mid D_p)$$

$$p(\bullet \mid D_1) = 4/5 = 0.8$$


$$p(\blacksquare \mid D_1) = 1/5 = 0.2 = 1 - p(\bullet \mid D_1)$$

# Decision trees – Sample probability at node

$p(c \mid D_n)$  : Sample probability of a sample having class  $c$  in data subset  $D_n$

$$\rightarrow p(c \mid D_n) = \frac{N_c}{\sum_{c \in K} N_c}$$

$$\sum_{c \in K} p(c \mid D_n) = 1$$



$$p(\bullet \mid D_p) = 5/11 = 0.45$$

$$p(\blacksquare \mid D_p) = 6/11 = 0.55 = 1 - p(\bullet \mid D_p)$$

$$p(\bullet \mid D_2) = 1/6 = 0.1\bar{6}$$

$$p(\blacksquare \mid D_2) = 5/6 = 0.8\bar{3} = 1 - p(\bullet \mid D_2)$$


## Decision trees – Impurity at node

$p(c \mid D_n)$  : Sample probability of a sample having class  $c$  in data subset  $D_n$

$$\rightarrow p(c \mid D_n) = \frac{N_c}{\sum_{c \in K} N_c}$$

Using classification error

$$I_E(D_n) = 1 - \max_{c \in K} \{ p(c \mid D_n) \}$$



# Decision trees – Information gain by split

## Using classification error


$$I_E(D_n) = 1 - \max_{c \in K} \{ p(c \mid D_n) \}$$

## Information gain

$$IG = I(D_p) - \sum_{j=1}^2 \frac{N_j}{N_p} I(D_j)$$

$$\begin{aligned} &= 0.45 - (5/11 \cdot 0.2) \\ &\quad - (6/11 \cdot 0.1\bar{6}) \\ &\approx 0.268 \end{aligned}$$

$$I_E(D_1) = 0.2$$



# Decision trees – Information gain by split 2 (hypothetical)

## Using classification error

$$I_E(D_n) = 1 - \max_{c \in K} \{ p(c \mid D_n) \}$$


## Information gain

$$IG = I(D_p) - \sum_{j=1}^2 \frac{N_j}{N_p} I(D_j)$$

$$= 0.45 - (5/11 \cdot 0) - (6/11 \cdot 0)$$

$$= 0.45$$

$$I_E(D_1) = 0$$



pure leaf nodes



# Decision trees – Impurity measures

Classification error: 
$$I_E(D_n) = 1 - \max_{c \in K} \{ p(c | D_n) \}$$

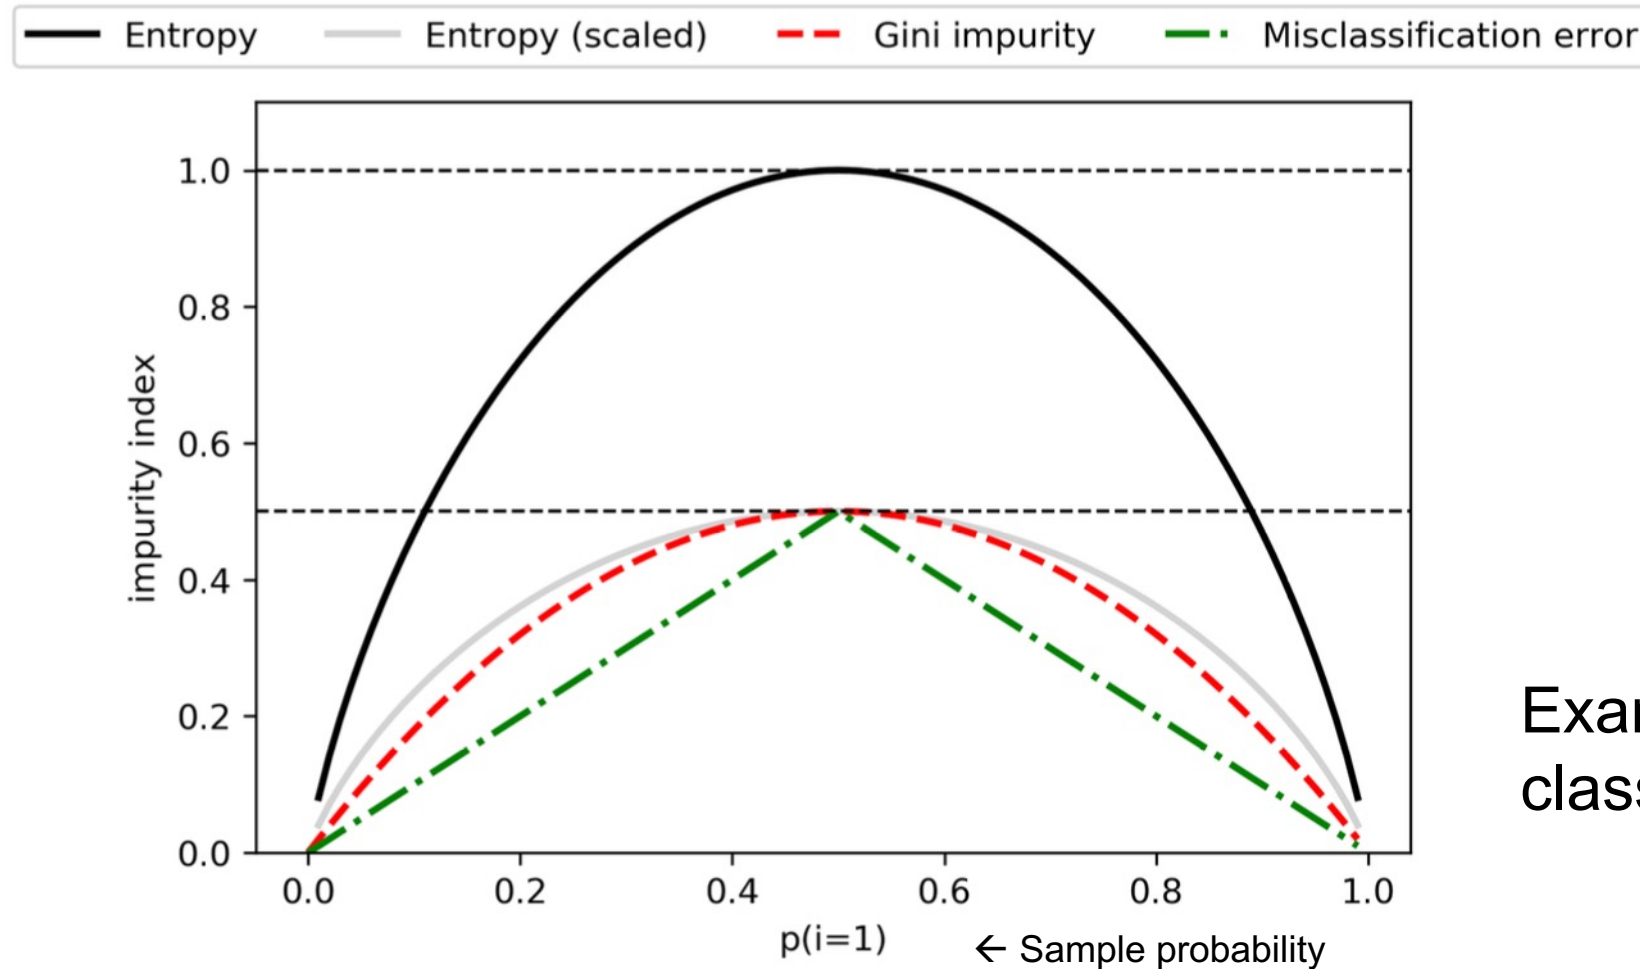
Entropy: 
$$I_H(D_n) := - \sum_{c \in K} p(c | D_n) \log_2 p(c | D_n)$$

(Information theory,  
Shannon 1948)

Gini impurity: 
$$I_{Gini}(D_n) = \sum_{c \in K} p(c | D_n) (1 - p(c | D_n)) = 1 - \sum_{c \in K} p(c | D_n)^2$$

**Impurity measure (I)** of a dataset associated with a node: a **measure** of how far away from a **pure** node (all samples one class) we are

# Decision trees – Impurity measures (overview)



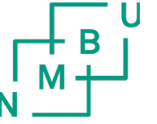
Example for binary classification



# Decision trees – Impurity measures (overview)

- In practice **Gini impurity** or **Entropy** are used
- They are **differentiable** functions (classification error is not) which simplifies finding optimal split
- They are more sensitive to changes in class probability and prefer splits which result in pure nodes with higher probability





# Decision trees – Scikit-learn example Iris

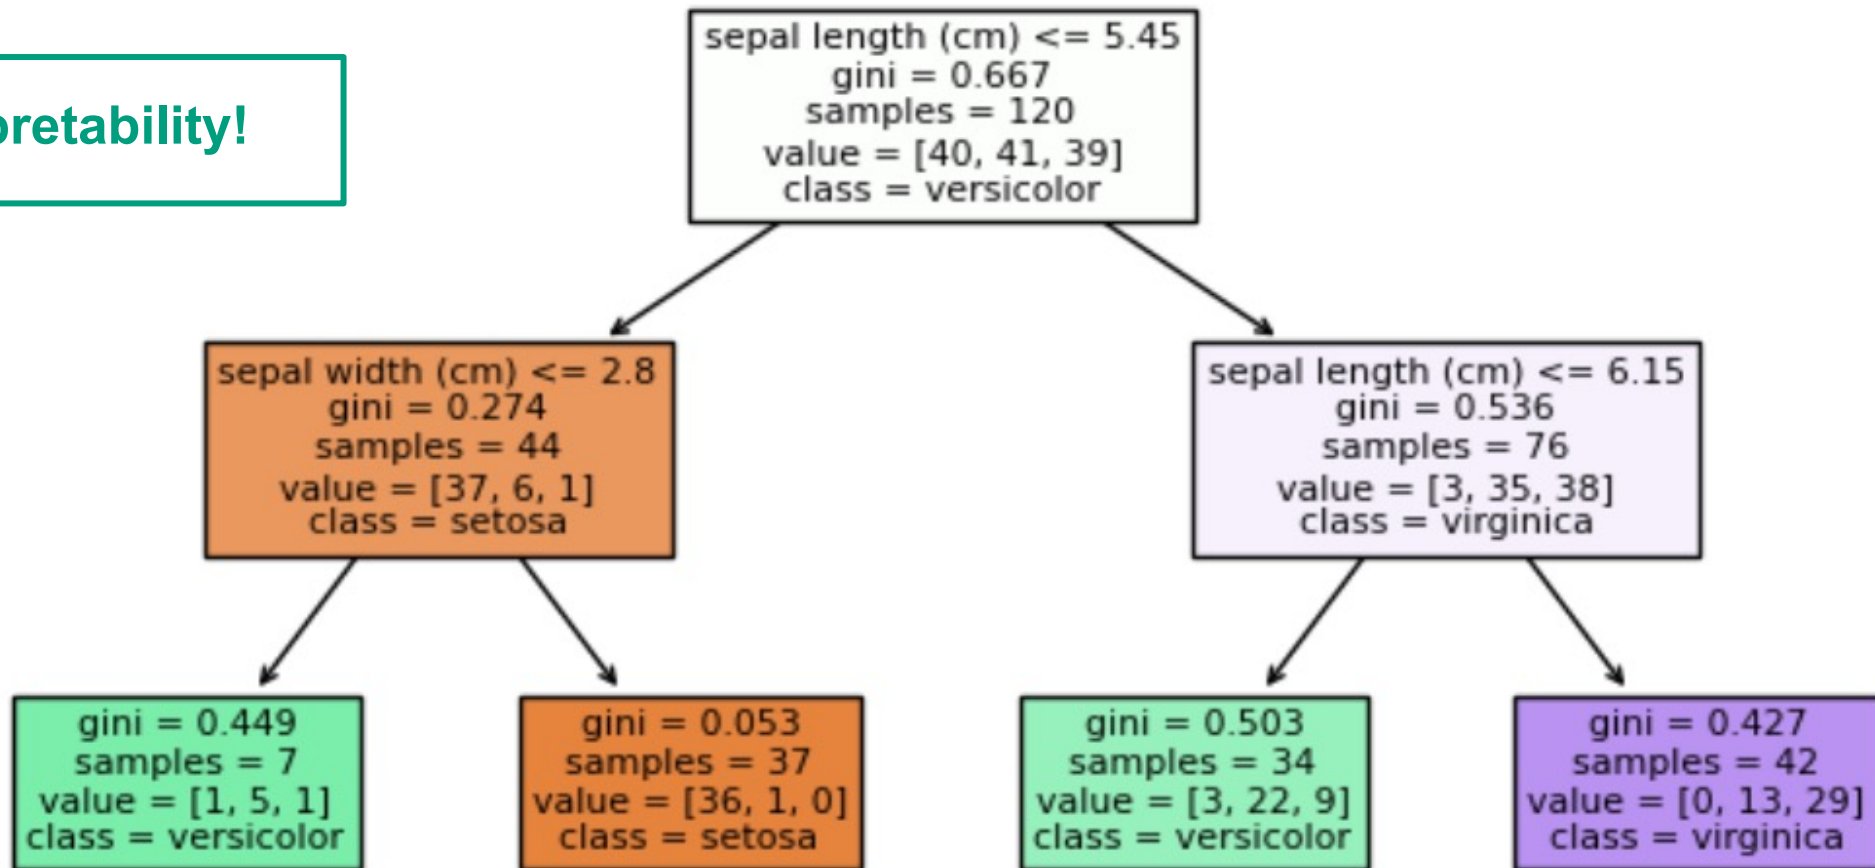
- Train a decision tree on the iris data set
- `from sklearn.tree import DecisionTreeClassifier`
- Plot the decision tree graphically
- `from sklearn.tree import plot_tree`

```
03_decision_tree_iris.ipynb
```

- Examples for hyperparameters: `max_depth`, `min_samples_leaf`, `criterion`
- <https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html>

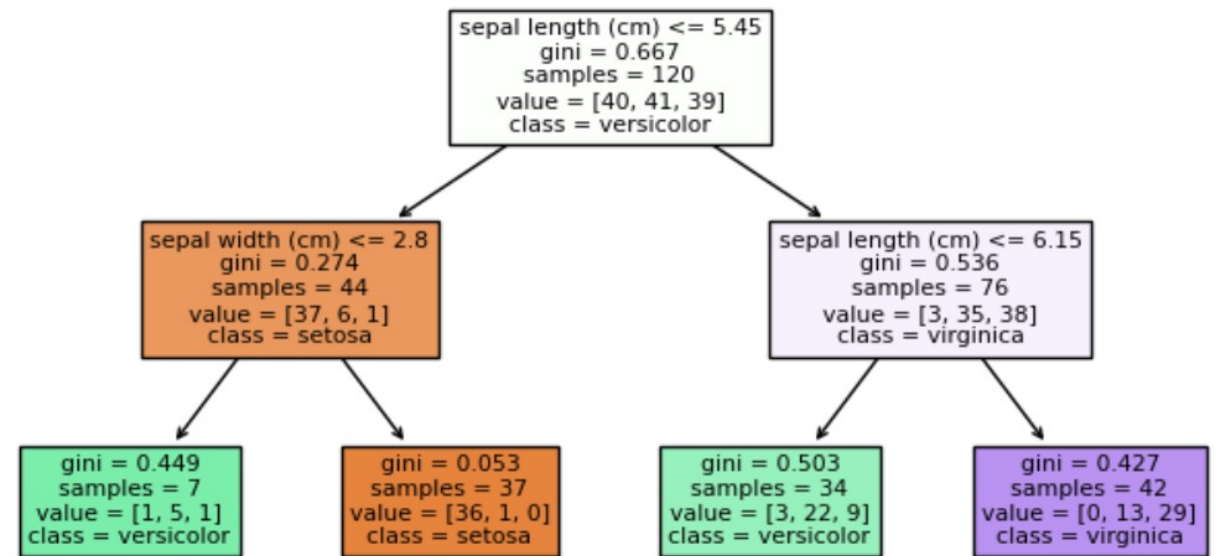

# Decision trees – Scikit-learn example Iris

Interpretability!



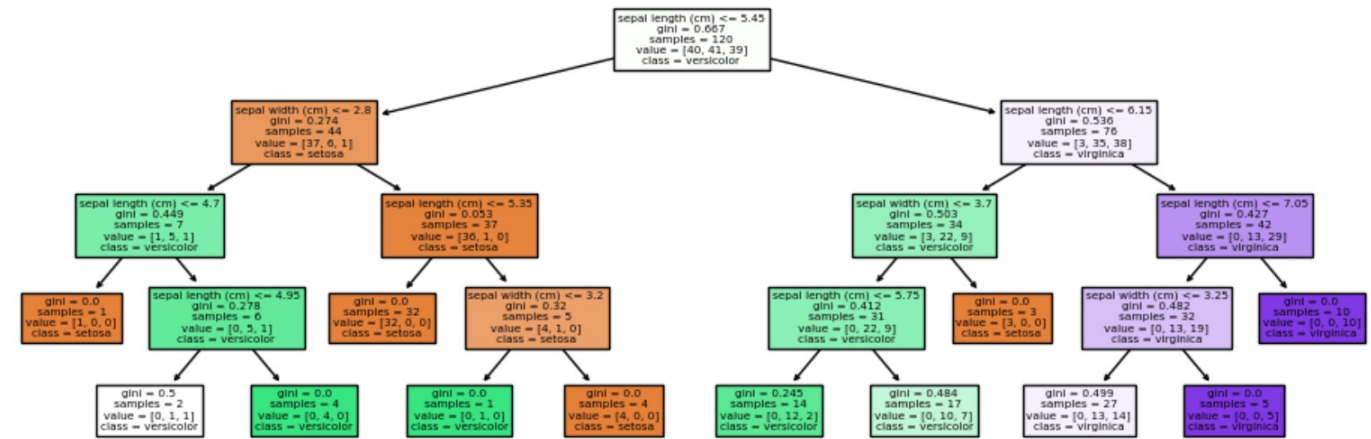
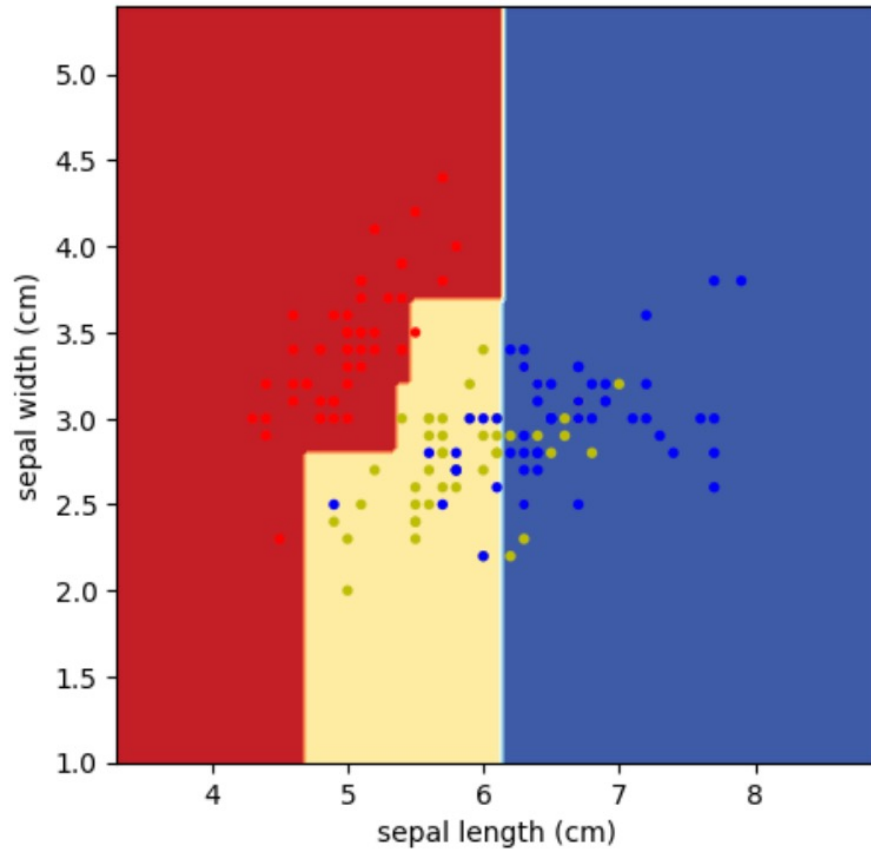
03\_decision\_tree\_iris.ipynb

# Decision trees – Scikit-learn example Iris (max\_depth=2)



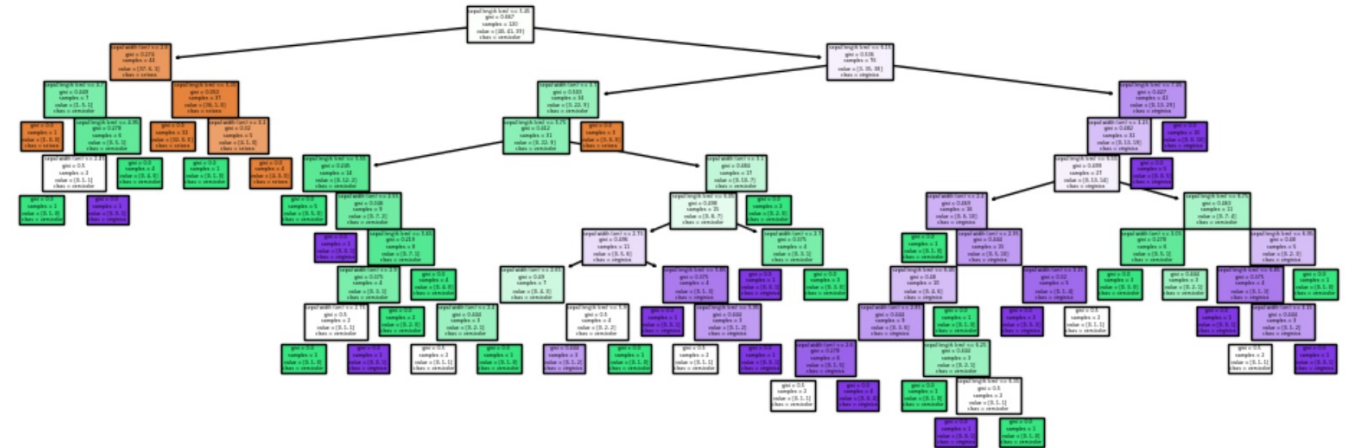
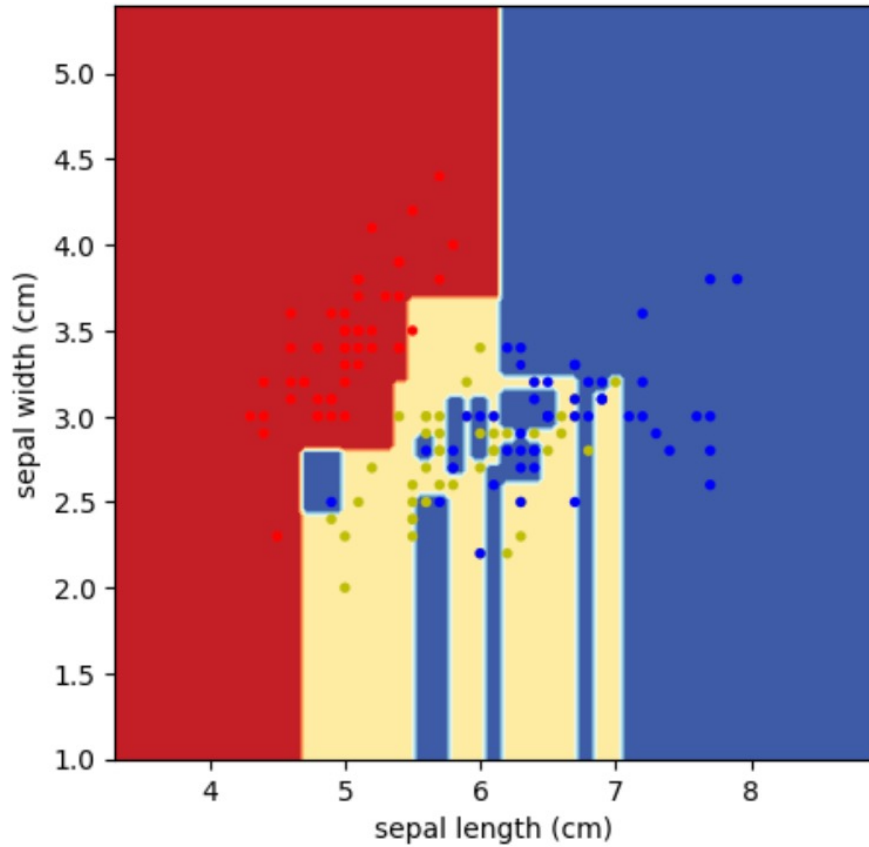
03\_decision\_tree\_iris.ipynb

# Decision trees – Scikit-learn example Iris (max\_depth=4)



03\_decision\_tree\_iris.ipynb

# Decision trees – Scikit-learn example Iris (max\_depth=20)



03\_decision\_tree\_iris.ipynb

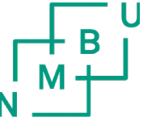
Overfitting!



# Decision trees – Pruning trees


- Tuning the **tree depth** or **minimal number of samples in a leaf node** e.g. via **grid search**
- Or, for example, pruning via “*cost-complexity analysis*” in sci-kit learn (not syllabus)

```
03_decision_tree_cost_complexity.ipynb
```



# Decision trees (summary)

- **Decision trees** can build **complex (non-linear)** decision boundaries by dividing the feature space into **rectangles**
- They are easy to interpret
- Need to be careful regarding **depth** of the decision tree
  - more complex the decision boundaries can easily result in **overfitting**
- **Note:** Feature scaling is **not** a requirement for decision tree algorithms
  - splits are easier to interpret with original scale



# Scikit-learn and Tour of Classifiers

## Random Forests





# Random forests

- Combining **multiple decision trees** into a powerful classifier
- **Bagging** and **boosting** are methods to combine trees (more after Easter break)
- Instance of an **ensemble learning** algorithm (more detailed after Easter break)
- **Average** multiple trees (each **high variance**) to obtain a more **robust** model
- Very popular due to
  - their good classification performance
  - their scalability
  - their ease of use



# Random forests – Bagging

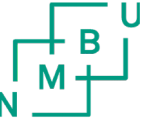
- **Goal:** Build a more **robust** model that has a **better generalisation** performance and is **less susceptible** to **overfitting**
- **Bagging** stands for **bootstrap aggregating** (Breiman 1994)
- Averaging **decreases** the **variance** of the model, **without increasing** the **bias**



# Random forests – Bagging

Random forest algorithm can be summarised in **four** simple steps

1. Draw a random **bootstrap sample** of size  $n$  (**randomly choose**  $n$  samples from the training set **with replacement**)
2. **Grow a decision tree** from the bootstrap sample. At each node
  - **Randomly select  $d$  features** (without replacement)
  - **Split** the node using the feature that provides the **best split** according to the objective function, for instance, **maximising the information gain**
3. **Repeat** the steps 1. and 2.  **$k$  times** ( $k$ : number of trees to be computed)
4. **Aggregate** the **prediction** by each tree to assign the class label by **majority vote**



# Random forests – Bagging in images (bootstrap sample)

Volume	Clusters	Shape factor	Malignant?
350	4	1	no
100	1	1	no
260	1	3.5	yes
10	3	2.5	yes

# Random forests – Bagging in images (bootstrap sample)

Volume	Clusters	Shape factor	Malignant?		Volume	Clusters	Shape factor	Malignant?
350	4	1	no	→	350	4	1	no
100	1	1	no					
260	1	3.5	yes					
10	3	2.5	yes					

**1. Random pick**

# Random forests – Bagging in images (bootstrap sample)

Volume	Clusters	Shape factor	Malignant?
350	4	1	no
100	1	1	no
260	1	3.5	yes
10	3	2.5	yes

Volume	Clusters	Shape factor	Malignant?
350	4	1	no
260	1	3.5	yes

2. Random pick

# Random forests – Bagging in images (bootstrap sample)

Volume	Clusters	Shape factor	Malignant?
350	4	1	no
100	1	1	no
260	1	3.5	yes
10	3	2.5	yes

Volume	Clusters	Shape factor	Malignant?
350	4	1	no
260	1	3.5	yes
100	1	1	no

3. Random pick

# Random forests – Bagging in images (bootstrap sample)

Volume	Clusters	Shape factor	Malignant?
350	4	1	no
100	1	1	no
260	1	3.5	yes
10	3	2.5	yes

Volume	Clusters	Shape factor	Malignant?
350	4	1	no
260	1	3.5	yes
100	1	1	no
100	1	1	no

4. Random pick



# Random forests – Bagging in images (bootstrap sample)

Original data

Volume	Clusters	Shape factor	Malignant?
350	4	1	no
100	1	1	no
260	1	3.5	yes
10	3	2.5	yes

Bootstrapped sample

Volume	Clusters	Shape factor	Malignant?
350	4	1	no
260	1	3.5	yes
100	1	1	no
100	1	1	no



Some samples don't occur



Sample can occur multiple times

# Random forests – Bagging in images (select features)

Volume	Clusters	Shape factor	Malignant?
350	4	1	no
100	1	1	no
260	1	3.5	yes
10	3	2.5	yes

Randomly select features




Volume	Shape factor	Malignant?
350	1	no
100	1	no
260	3.5	yes
10	2.5	yes

Compute split that maximizes information gain among the selected features


Repeat for every node

Each split at any node may use a different subset of features



# Random forests – Bagging in images (many trees)




## Train many trees – each will look a bit different



# Random forests – Bagging in images (voting)

Train many trees – each will look a bit different

New data >

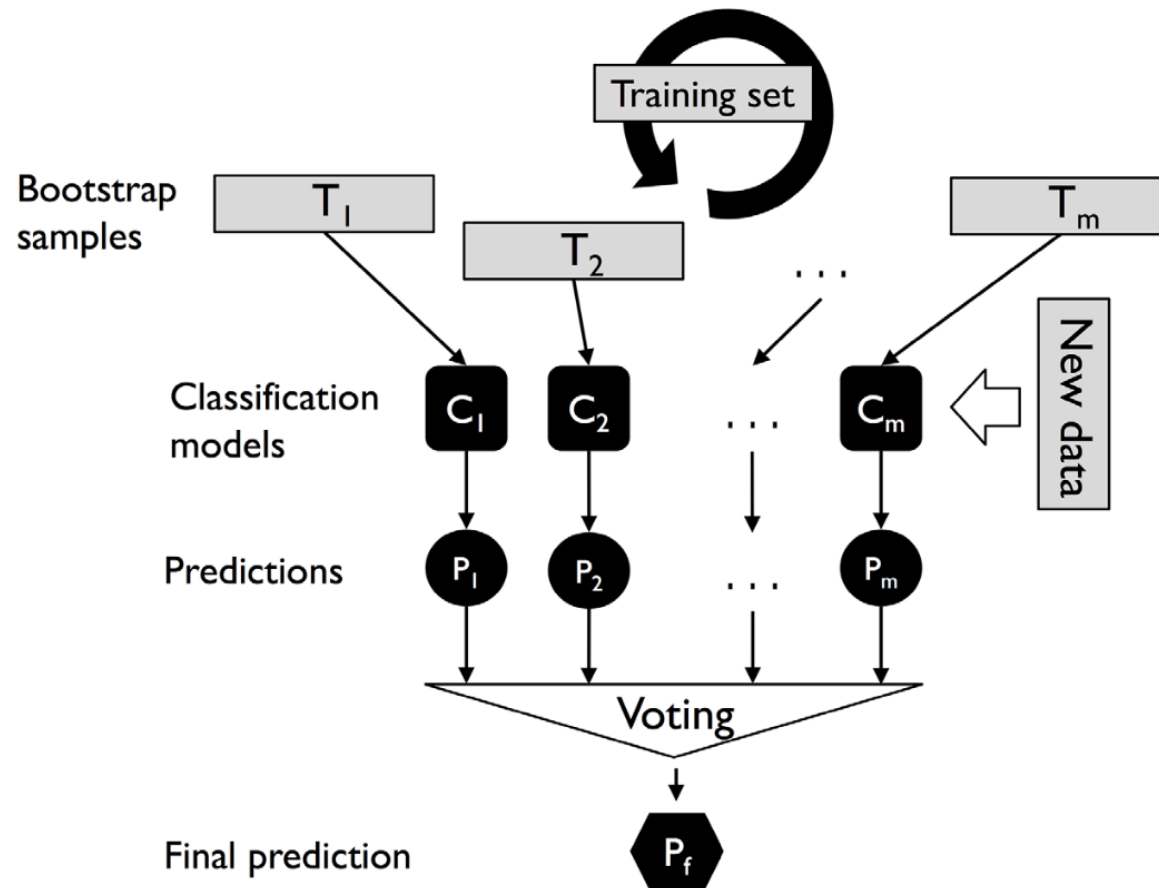


...


predict  
74 yes / 26 no

Majority vote: **yes**

# Random forests – Bagging summary



# Random forest classifier (code example iris)



max\_depth=3  
n\_estimators=100  
n\_jobs=2 (compute tree in parallel)

**Decision boundary is smoother because it's obtained by a majority vote over 100 high variance decision trees**

**Random forest has usually a better generalization error than a single tree**

`03_random_forest_iris.ipynb`

`03_randomforest_and_decisontree.ipynb`



## Random forest classifier (summary)

- Robust classifier by averaging over many trees
- Do not offer the same level of interpretability as plain decision trees
- Typically, the more trees are used, the better
- But: using more than one tree results in higher computational effort
- Less hyperparameters to tune due to robustness
  - Number of trees (`n_estimators`)
  - Size of bootstrap samples → usually fixed to training size
  - Number of features to randomly select → usually fixed to  $\sqrt{m}$  for  $m$  features

