


# Hyperparameter optimization


Grid-search, random-search and other methods for hyperparameter-search

see Ch. 05 in book “Python Machine Learning” by Raschka & Mirjalili



# Overview

- Binary classifier example
- Confusion matrix
- Different evaluation metrics
- Receiver Operator Curve (ROC) and Area Under Curve (ROC-AUC)
- Performance evaluation for multiclass problems
  - ⋈ Micro-averaging
  - ⋈ Macro-averaging
- Other metrics worth knowing about



## Binary classifier example

- Rare type of asthma occurs in approximately 1% of the population
- Can be deadly, but is a condition that is easy to treat
- Assume that you have two binary classifiers:
  - ⋈ C1 with an accuracy of 99%
  - ⋈ C2 with an accuracy of 98%
- Which one would you prefer?

# Binary classifier example

- Let us take a close look on how the two classifiers perform on the test set
- 1000 samples in total 990 healthy and 10 with the disease
- C1 makes 10 mistakes, C2 makes 18 mistakes
- Does anyone see any issue with C1, despite it being more accurate than C2?

IDX	1	2	3	...	973	974	975	976	977	978	979	980	981	982	983	984		
GT	0	0	0	...	0	0	0	0	0	0	0	0	0	0	0	0	No disease	Disease
C1	0	0	0	...	0	0	0	0	0	0	0	0	0	0	0	0	Correct	Incorrect
C2	0	0	0	...	0	1	1	1	1	1	1	1	1	1	1	1	Correct	Incorrect
IDX	985	986	987	988	989	990	991	992	993	994	995	996	997	998	999	1000		
GT	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	No disease	Disease
C1	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	1	Correct	Incorrect
C2	1	1	1	1	1	1	0	1	1	1	1	1	1	1	1	1	Correct	Incorrect

# Confusion matrix

- Often we want to get a better idea of how well a classifier detects specific classes, and an overall accuracy is not enough
- To get a more thorough understanding of a models predicted performance we look at a models **confusion matrix**
- For binary classifiers it looks like to matrix to the right

		Predicted class	
		$P$	$N$
Actual class	$P$	True positives (TP)	False negatives (FN)
	$N$	False positives (FP)	True negatives (TN)

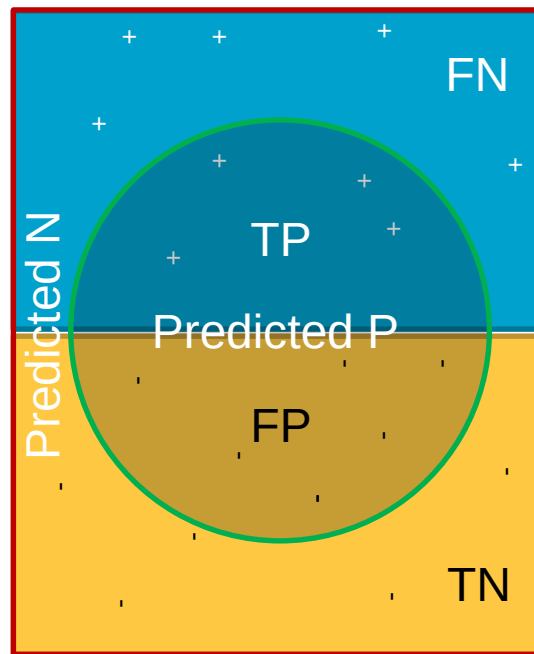
# Confusion matrix

- It is a square matrix showing the classes that a model predicts versus the classes of the ground truth
- Let 1 be positive and 0 be negative
- **TP**: Positive sample model predicts positive
- **FP**: Negative sample model predicts positive
- **FN**: Positive sample model predicts negative
- **TN**: Negative sample model predicts negative


		Predicted class	
		$P$	$N$
Actual class	$P$	True positives (TP)	False negatives (FN)
	$N$	False positives (FP)	True negatives (TN)

# Confusion matrix

Bad model

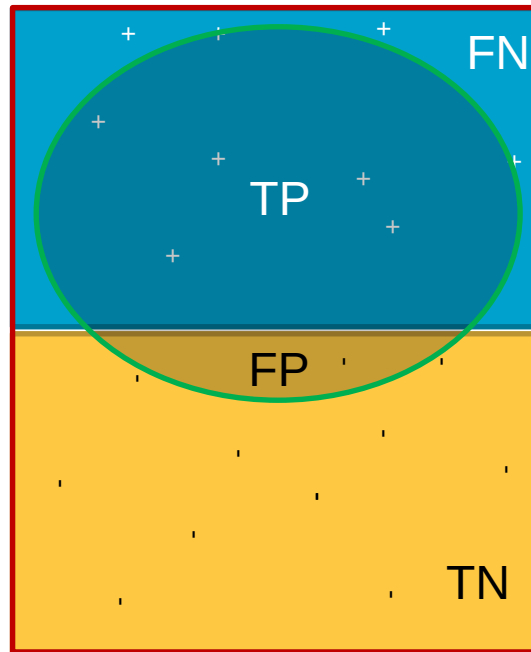


Better model




# Confusion matrix

Balanced classes



Unbalanced classes



Predict  
everything  
positive?






# Confusion matrix

- Lets compute the confusion matrix of CA1 and CA2 from the rare asthma condition example

			<i>PP</i>	<i>PN</i>			
		<i>P</i>	TP	FN			
		<i>N</i>	FP	TN			
	<b>C1</b>				<b>C2</b>		
			<i>PP</i>	<i>PN</i>			
<i>P</i>		1	9		<i>P</i>	9	1
<i>N</i>		1	989		<i>N</i>	17	973



# Different evaluation metrics

- As of now you should be familiar with two evaluation metrics
  - ⋈ **Accuracy (ACC)**
  - ⋈ **Prediction error (ERR)**
- **ACC**: Number of **correct** predictions divided by number of total predictions
- **ERR**: Number of **incorrect** predictions divided by number of total predictions

# Different evaluation metrics

- As of now you should be familiar with two evaluation metrics
  - ↪ **Accuracy (ACC)**
  - ↪ **Prediction error (ERR)**

$$\mathbf{ACC} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\mathbf{ERR} = \frac{FP + FN}{TP + TN + FP + FN} = 1 - \mathbf{ACC}$$



# Different evaluation metrics

- Different evaluation metrics are important for different classification problems
- True Positive Rate (**TPR**), False Positive Rate (**FPR**), False Negative Rate (**FNR**) and True Negative Rate (**TNR**) are especially useful for imbalanced datasets
- In some problems it will be important to minimize the FNR, in other problems it will not be as important

$$\mathbf{TPR} = \frac{TP}{P} = \frac{TP}{TP + FN}$$

$$\mathbf{FPR} = \frac{FP}{N} = \frac{FP}{FP + TN}$$

$$\mathbf{FNR} = \frac{FN}{P} = \frac{FN}{FN + TP}$$

$$\mathbf{TNR} = \frac{TN}{N} = \frac{TN}{TN + FP}$$


## Different evaluation metrics

- **Precision** (PRE), **Recall** (REC) and **F1**-score are central metrics in this course
- Precision seeks to measure the amount of TP's in relation to FP's
- Recall (also known as *sensitivity* in medicine) is equivalent to TPR

$$\text{Recall} = \frac{TP}{TP + FN} \quad \text{Precision} = \frac{TP}{TP + FP}$$

- Often, optimizing for recall might come at the cost of lowering precision
- F1-score is a metric that seeks to combine precision and recall

$$\text{F1} = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$




# Examples precision recall and F1

- Let us compute the precision, recall and F1-scores of our C1 and C2 classifiers for the Asthma example

# Receiver Operator Curve and Area Under Curve


- The Receiver Operator Curve (ROC) is a graphical representation of a classifier performance.
- Plots TPR versus FPR for a binary classifier at **different decision thresholds**
- Can be understood as trying to visualize how much better a model is than random guessing
- Can be used with any classifier that applies a decision boundary (the majority of classifiers)
- [Online illustration](#)





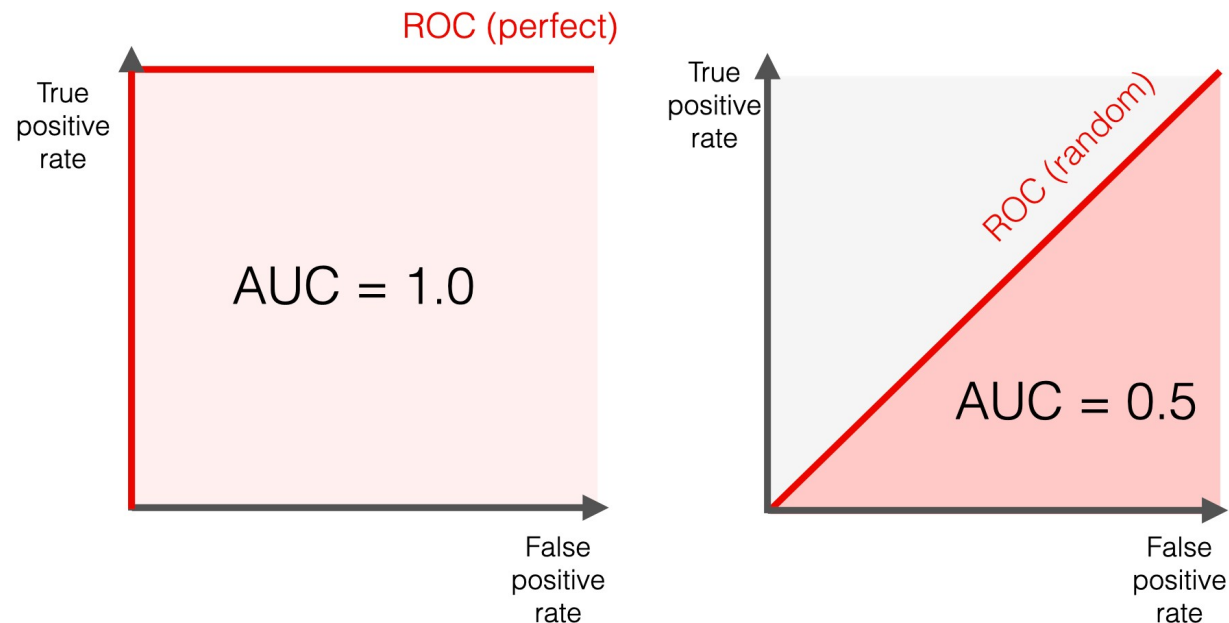
# Receiver Operator Curve and Area Under Curve

- Illustrates the trade-off between FPR and TPR
- To compute the ROC classifiers must output a probabilistic output/softmax output which can be thresholded
- Diagonal curve is the performance of random guessing (worst possible score)
- Blue curve is better



# Receiver Operator Curve and Area Under Curve

- ROC-AUC: Area Under the Receiver Operating Characteristic Curve
- Quantitative measure of overall classifier performance at all possible thresholds



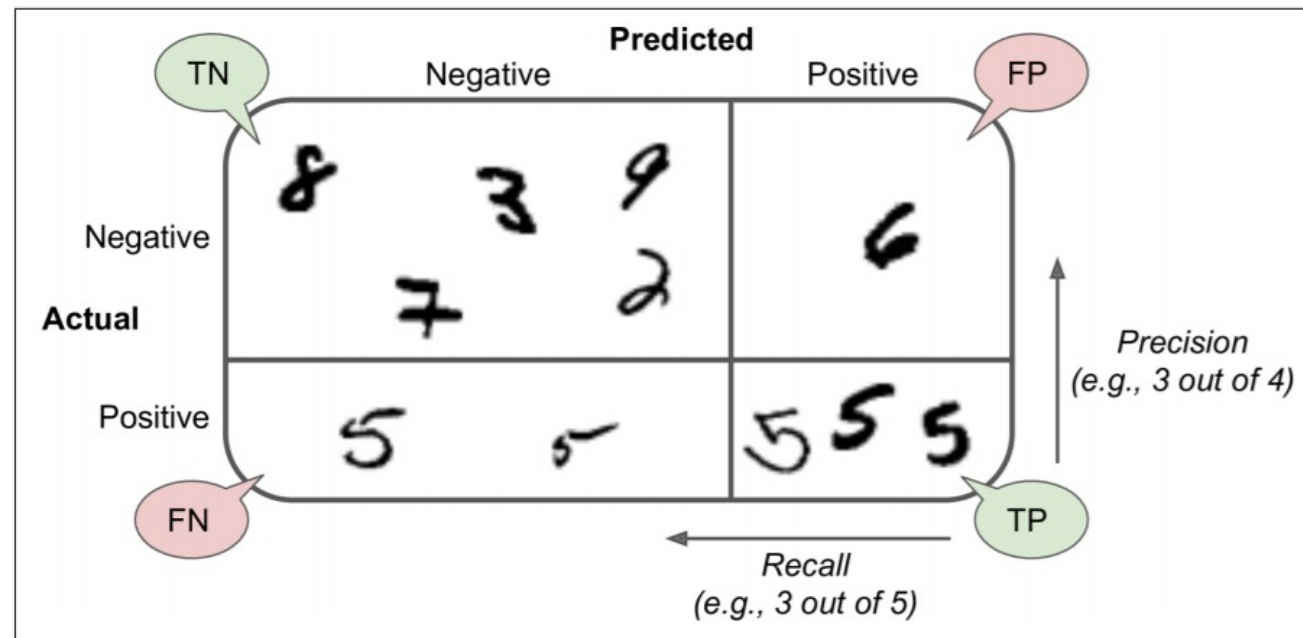



# Examples ROC-AUC

- `ROC_AUC_with_LR.ipynb`
  - ⌘ Example of how visualize the ROC and compute the AUC

# Performance evaluation for multiclass problems

- When evaluating multiclass models: compute multiple “one-vs-all” assessments
- This is what one one-vs-all assessment looks like for handwritten digits: 5 vs. “not 5”



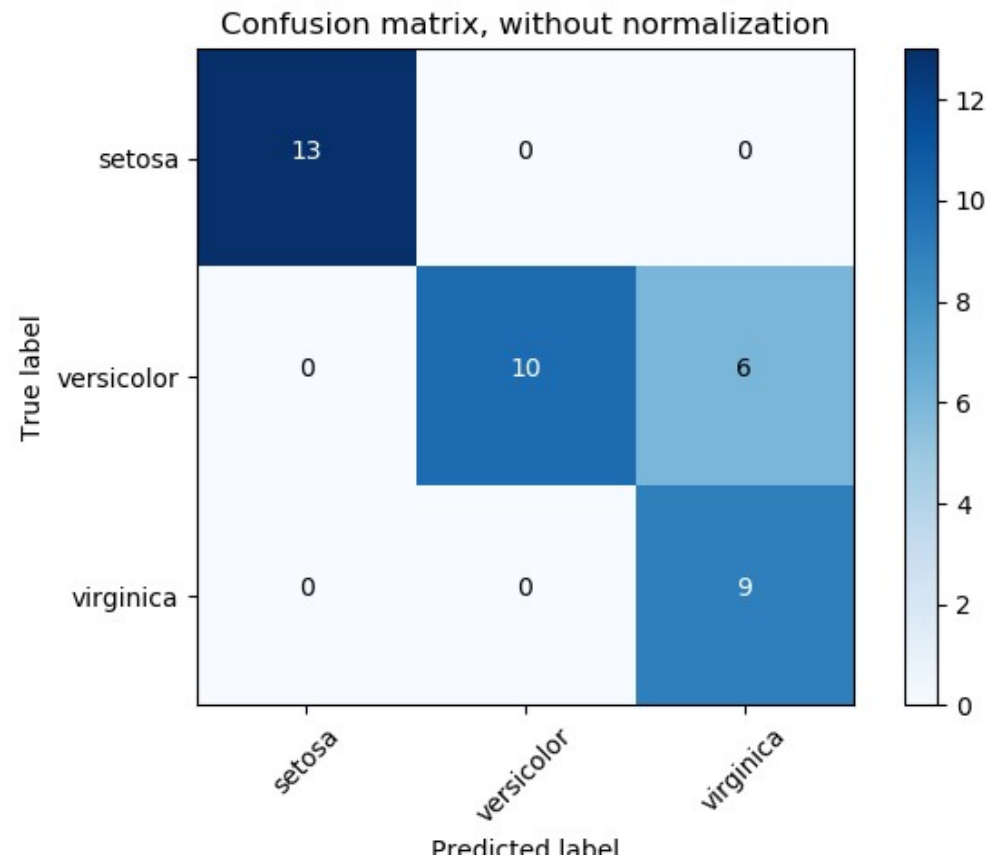


# Performance evaluation for multiclass problems

- Can choose between **micro** and **macro** averaging:
  - Micro gives **equal weight** to every **sample**
  - Macro gives **equal weight** to every **class** (sample/class\_size)
- Micro/macro averaging can be done for any performance metric
- Can even be used for binary classification problems

# Performance evaluation for multiclass problems


- Confusion matrices for  $n$ -class multiclass problems have  $n \times n$  dimensions



		Predicted class	
		$P$	$N$
Actual class	$P$	True positives (TP)	False negatives (FN)
	$N$	False positives (FP)	True negatives (TN)

# Performance evaluation for multiclass problems

- Let us use this confusion matrix to illustrate difference between micro/macro averaging




# Performance evaluation for multiclass problems

- Lets look at the example from the beginning of the class to show how micro/macro averaging can be used for binary classifiers


			<i>PP</i>	<i>PN</i>			
		<i>P</i>	TP	FN			
		<i>N</i>	FP	TN			
	<b>C1</b>				<b>C2</b>		
			<i>PP</i>	<i>PN</i>			
	<i>P</i>	1	9		<i>P</i>	9	1
	<i>N</i>	1	989		<i>N</i>	17	973





# Performance evaluation for multiclass problems

- There are almost as many evaluation metrics as there are problems
- Different people also use different names for the same metric
  - TPR/sensitivity/recall
  - precision/positive predictive value
- Some are application specific
  - Computer vision: mean Average Precision, Intersection over Union, etc.
  - Medicine: Diagnostic Odds Ratio
  - Unsupervised learning: Adjusted Rand Index



# Performance evaluation for multiclass problems


- Some of my favourites include:
  - F1-score
  - Matthews Correlation Coefficient (MCC)
  - Balanced Accuracy

# Performance evaluation for multiclass problems

Check out the [Wikipedia page](#), or [scikit-learn documentation](#) for a good overview


Sources: [4][5][6][7][8][9][10][11][12] view · talk · edit

		Predicted condition			
		Total population = P + N	Predicted Positive (PP)	Predicted Negative (PN)	
Actual condition	Positive (P) [a]	True positive (TP), hit <sup>[b]</sup>	False negative (FN), miss, underestimation	True positive rate (TPR), recall, sensitivity (SEN), probability of detection, hit rate, power $= \frac{TP}{P} = 1 - FNR$	Prevalence threshold (PT) $= \frac{\sqrt{TPR \times FPR} - FPR}{TPR - FPR}$
	Negative (N) <sup>[d]</sup>	False positive (FP), false alarm, overestimation	True negative (TN), correct rejection <sup>[e]</sup>	False positive rate (FPR), probability of false alarm, fall-out type I error <sup>[f]</sup> $= \frac{FP}{N} = 1 - TNR$	False negative rate (FNR), miss rate type II error <sup>[c]</sup> $= \frac{FN}{P} = 1 - TPR$
	Prevalence $= \frac{P}{P + N}$	Positive predictive value (PPV), precision $= \frac{TP}{PP} = 1 - FDR$	False omission rate (FOR) $= \frac{FN}{PN} = 1 - NPV$	Positive likelihood ratio (LR+) $= \frac{TPR}{FPR}$	Negative likelihood ratio (LR-) $= \frac{FNR}{TNR}$
	Accuracy (ACC) $= \frac{TP + TN}{P + N}$	False discovery rate (FDR) $= \frac{FP}{PP} = 1 - PPV$	Negative predictive value (NPV) $= \frac{TN}{PN} = 1 - FOR$	Markedness (MK), deltaP ( $\Delta p$ ) $= PPV + NPV - 1$	Diagnostic odds ratio (DOR) $= \frac{LR+}{LR-}$
	Balanced accuracy (BA) $= \frac{TPR + TNR}{2}$	F <sub>1</sub> score $= \frac{2 PPV \times TPR}{PPV + TPR} = \frac{2 TP}{2 TP + FP + FN}$	Fowlkes-Mallows index (FM) $= \sqrt{PPV \times TPR}$	Matthews correlation coefficient (MCC) $= \frac{\sqrt{TPR \times TNR \times PPV \times NPV}}{\sqrt{FNR \times FPR \times FOR \times FDR}}$	Threat score (TS), critical success index (CSI), Jaccard index $= \frac{TP}{TP + FN + FP}$




## Other metrics worth knowing about

- You don't need to memorize the formulas for all the evaluation metrics that you have been presented in this lecture for the exam
- Understand the **concepts** of the metrics such that you can use them later
  - ⋈ Should be able to remember what the purpose of a confusion matrix
  - ⋈ Assess when it is appropriate to use which metric
  - ⋈ When working with multiclass, when one should use micro/macro averaging



## Other metrics worth knowing about

- That being said, there are a few things you should memorize for the exam
  - ↪ Confusion matrix
  - ↪ Precision, Recall and F1-score
  - ↪ Concept of an ROC curve, and ROC-AUC metric
  - ↪ Difference between micro and macro averaging when computing multiclass metrics



# Examples metrics with scikit-learn

- `metrics_example.ipynb`
  - ⌘ Example of how compute evaluation metrics with scikit-learn

Thank you for listening

